

Getting Under and Over Vegas NFL Lines

```
In [632]: from IPython.display import Image
Image("/Users/matthewnykaza/Documents/Flatiron/Phase_3_Project/NFL_Betting_Data/images/vegas_image.jpeg")
```

```
Out[632]:
```



Day	Team	Line	
FRIDAY OCTOBER 6	FLA ST	-11	
	NC ST	-1	
	TCU	-1	
	UTAH	-1	
	LOUISVILLE	-33	
	MD TN ST*	-33	
	SATURDAY OCTOBER 7	W WESTERN	-20.5
		WISCONSIN	-20.5
		URDUE	-11
		OWA	-11
TTSBURG		-6.5	
ACUSE		-6.5	
DIANA		-7	
INOIS		-7	
MSON		-16	
FOREST		-16	
AR	-16		
ALABAMA*	-29		
OKLA ST	-3		
KANSAS ST	-3		
TEXAS A&M	-1.5		
KANSAS	-27		
S DIEGO S	-27		
BYU	-2		
RICE	-2		
TULANE	-3		
NAVY	-3		
AIR FORCE	-32		
STANFORD	-32		
N TRE DAME	-28		
W VIRGINI	-28		
MISS ST	-2		
LSU	-2		
FLORIDA	-4		
WASH ST	-4		
OREGON ST	-1.5		
ARIZONA	-1.5		
UCLA	-1.5		
VIRGINIA	-6.5		
E CAROLNA	-6.5		
AKRON	-8		
CINCINATI	-8		
S CAR	-8.5		
KENTUCKY	-8.5		
BAYLOR	-5		
COLORADO	-5		
NEBRASKA	-6.5		
IOWA ST	-8		
ALA-BIRM	-8		
MISSOURI	-3.5		
TEX TECH	-3.5		
WEST MICH	-3		
OHIO	-1.5		
VANDY	-1.5		
OLE MISS	-1.5		
CENT MICH	-1.5		
TOLSON	-1.5		

Overview

For this project I sourced NFL betting, stadiums and teams data from the data website Kaggle. This data includes data going back to the 1978 NFL seasons creating a classification model that will assist sports bettors with selecting the over or under in an NFL game. The Over/Under line is the predicted combined the two teams must combine for greater than that total, and to go Under they must combine for less than that total. I began this project by preprocessing the variables, engineering relevant ones from the given data, and performing intelligent decisions as what to do with missing/incomplete data. Once the processing using that data, and Sklearn's pipeline function. After some early attempts with base a Logistic Regression model, it was determined that using a tree-based scores. The final model used was a Random Forest model, with some hyperparameter tuning, that allowed me to achieve roughly 53% accuracy on my test data. I will continue to tune the model to achieve optimal results.

```
In [603]: import numpy as np
import pandas as pd
import requests
from geopy.extra.rate_limiter import RateLimiter
from geopy.geocoders import Nominatim
import geocoder
import matplotlib.pyplot as plt
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score, RandomizedSearchCV
from sklearn.metrics import accuracy_score, f1_score, roc_auc_score, plot_confusion_matrix, roc_curve, auc
from sklearn.neighbors import KNeighborsClassifier
import category_encoders as ce
import seaborn as sns
from statsmodels.stats.outliers_influence import variance_inflation_factor
import xgboost as xgb
```

The Datasets

For this project I had three datasets

- NFL betting information which included dates, week of schedule, teams, scores, betting lines, weather, stadium names, and playoff information
- NFL stadium data which had more detailed information about the various stadiums that NFL teams have played at since 1978
- NFL team data which included information about individual teams such as nicknames, conference and division information The main data that I used was information to dig into greater detail about individual stadiums, and I used the NFL teams data to compare conference/divisional matchups, as well as each team's last 5 games.

The main data that I used was the NFL betting, but I used the stadium information to dig into greater detail about individual stadiums, and I used the NFL team matchups, as well as help setup the average scores for each individual team's last 5 games.

Load in Datasets

```
In [ ]:

In [2]: nfl_teams = pd.read_csv("//Users/matthewnykaza/Documents/Flatiron/Phase_3_Project/NFL_Betting_Data/Files/nfl_teams.csv")

In [171]: nfl_teams2 = pd.read_csv("//Users/matthewnykaza/Documents/Flatiron/Phase_3_Project/NFL_Betting_Data/Files/nfl_teams.csv")

In [3]: nfl_teams.loc[nfl_teams['team_name'] == 'Las Vegas Raiders']
```

Out[3]:

	team_name	team_name_short	team_id	team_id_pfr	team_conference	team_division	team_conference_pre2002	team_division_pre2002
31	Las Vegas Raiders	Raiders	LVR	RAI	AFC	NaN	AFC	AFC West

```
In [4]: nfl_teams.loc[nfl_teams['team_name'] == 'Las Vegas Raiders', 'team_division'] = 'AFC West'
```

```
In [5]: nfl_stadiums = pd.read_csv("//Users/matthewnykaza/Documents/Flatiron/Phase_3_Project/NFL_Betting_Data/Files/nfl_stadiums.csv")
nfl_stadiums.head()
```

Out[5]:

	stadium_name	stadium_location	stadium_open	stadium_close	stadium_type	stadium_address	stadium_weather_station_code	stadium_weather_type	stadium_weather_season
0	Alamo Dome	San Antonio, TX	NaN	NaN	indoor	100 Montana St, San Antonio, TX 78203	78203	dome	indoor
1	Allegiant Stadium	Paradise, NV	2020.0	NaN	indoor	NaN	NaN	dome	indoor
2	Altel Stadium	Jacksonville, FL	NaN	NaN	NaN	NaN	NaN	NaN	NaN
3	Alumni Stadium	Chestnut Hill, MA	NaN	NaN	outdoor	Perimeter Rd, Chestnut Hill, MA 02467	2467	cold	outdoor
4	Anaheim Stadium	Anaheim, CA	1980.0	1994.0	outdoor	2000 E Gene Autry Way, Anaheim, CA 92806	92806	warm	outdoor

```
In [6]: nfl_scores = pd.read_csv("//Users/matthewnykaza/Documents/Flatiron/Phase_3_Project/NFL_Betting_Data/Files/spreadspoke.csv")
nfl_scores.head()
```

Out[6]:

season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite	over_under_line	stadium	stadium_location
1966	1	False	Miami Dolphins	14.0	23.0	Oakland Raiders	NaN	NaN	NaN	Orange Bowl	Orlando, FL
1966	1	False	Houston Oilers	45.0	7.0	Denver Broncos	NaN	NaN	NaN	Rice Stadium	Houston, TX
1966	1	False	San Diego Chargers	27.0	7.0	Buffalo Bills	NaN	NaN	NaN	Balboa Stadium	San Diego, CA
1966	2	False	Miami Dolphins	14.0	19.0	New York Jets	NaN	NaN	NaN	Orange Bowl	Orlando, FL
1966	1	False	Green Bay Packers	24.0	3.0	Baltimore Colts	NaN	NaN	NaN	Lambeau Field	Green Bay, WI

Focus on NFL Stadiums Null Values

```
In [7]: nfl_stadiums.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 106 entries, 0 to 105
Data columns (total 15 columns):
 stadium_name          106 non-null object
 stadium_location      106 non-null object
 stadium_open          82 non-null float64
 stadium_close         41 non-null float64
 stadium_type          99 non-null object
 stadium_address       94 non-null object
 stadium_weather_station_code 93 non-null object
 stadium_weather_type  99 non-null object
 stadium_capacity      45 non-null object
 stadium_surface       59 non-null object
 STATION              55 non-null object
 NAME                 55 non-null object
 LATITUDE             55 non-null float64
 LONGITUDE            55 non-null float64
 ELEVATION             55 non-null float64
dtypes: float64(5), object(10)
memory usage: 12.5+ KB
```

A few quick thoughts

- We want this to help make sure we have all the data possible for a merge on nfl_scores
- Knowing indoor and outdoor stadiums will be important later on, and we can find that info on the web and do a .loc to update the info
- Once we have the address and stadium type's updated, we can import all latitudes, longitudes and elevation information
- We should also be able to get weather on a given day with this info aswell

```
In [8]: nfl_stadiums[nfl_stadiums['stadium_address'].isnull()]
```

Out[8]:

	stadium_name	stadium_location	stadium_open	stadium_close	stadium_type	stadium_address	stadium_weather_station_code	stadium_weather_type	stadium_elevation
1	Allegiant Stadium	Paradise, NV	2020.0	NaN	indoor	NaN	NaN	dome	1000
2	Alltel Stadium	Jacksonville, FL	NaN	NaN	NaN	NaN	NaN	NaN	100
18	Dolphin Stadium	Miami, FL	NaN	NaN	NaN	NaN	NaN	NaN	100
38	Jack Murphy Stadium	San Diego, CA	NaN	NaN	NaN	NaN	NaN	NaN	100
39	Joe Robbie Stadium	Miami, FL	NaN	NaN	NaN	NaN	NaN	NaN	100
55	Mercedes-Benz Stadium	Atlanta, GA	2017.0	NaN	indoor	NaN	NaN	dome	1000
69	Pro Player Stadium	Miami, FL	NaN	NaN	NaN	NaN	NaN	NaN	100

```
In [9]: nfl_stadiums[nfl_stadiums['stadium_type'].isnull()]
```

Out[9]:

	stadium_name	stadium_location	stadium_open	stadium_close	stadium_type	stadium_address	stadium_weather_station_code	stadium_weather_type	stadium_weather_type
2	Alltel Stadium	Jacksonville, FL	NaN	NaN	NaN	NaN	NaN	NaN	NaN
18	Dolphin Stadium	Miami, FL	NaN	NaN	NaN	NaN	NaN	NaN	NaN
38	Jack Murphy Stadium	San Diego, CA	NaN	NaN	NaN	NaN	NaN	NaN	NaN
39	Joe Robbie Stadium	Miami, FL	NaN	NaN	NaN	NaN	NaN	NaN	NaN
69	Pro Player Stadium	Miami, FL	NaN	NaN	NaN	NaN	NaN	NaN	NaN
84	Stanford Stadium	Palo Alto, CA	NaN	NaN	NaN	NaN	NaN	NaN	NaN
89	Tampa Stadium	Tampa, FL	NaN	NaN	NaN	NaN	NaN	NaN	NaN

```
In [10]: addresses = ['3333 Al Davis Way, Las Vegas, NV 89118', '1 TIAA Bank Field Dr, Jacksonville, FL 32202', '347 Don Shula Dr, Jacksonville, FL 32202']
types = ['outdoor', 'outdoor', 'outdoor', 'outdoor', 'outdoor', 'outdoor', 'outdoor']
```

```
In [11]: nfl_stadiums.loc[nfl_stadiums['stadium_address'].isnull(), 'stadium_address'] = addresses
nfl_stadiums.loc[nfl_stadiums['stadium_type'].isnull(), 'stadium_type'] = types
nfl_stadiums.head()
```

Out[11]:

	stadium_name	stadium_location	stadium_open	stadium_close	stadium_type	stadium_address	stadium_weather_station_code	stadium_weather_type	stadium_weather_type
0	Alamo Dome	San Antonio, TX	NaN	NaN	indoor	100 Montana St, San Antonio, TX 78203	78203	dome	dome
1	Allegiant Stadium	Paradise, NV	2020.0	NaN	indoor	3333 Al Davis Way, Las Vegas, NV 89118	NaN	dome	dome
2	Alltel Stadium	Jacksonville, FL	NaN	NaN	outdoor	1 TIAA Bank Field Dr, Jacksonville, FL 32202	NaN	NaN	NaN
3	Alumni Stadium	Chestnut Hill, MA	NaN	NaN	outdoor	Perimeter Rd, Chestnut Hill, MA 02467	2467	cold	cold
4	Anaheim Stadium	Anaheim, CA	1980.0	1994.0	outdoor	2000 E Gene Autry Way, Anaheim, CA 92806	92806	warm	warm

```
In [12]: nfl_stadiums['stadium_address'].isnull().sum()
```

Out[12]: 0

```
In [13]: nfl_stadiums['stadium_type'].isnull().sum()
```

```
Out[13]: 0
```

```
In [14]: nfl_stadiums.head(50)
```

```
Out[14]:
```

	stadium_name	stadium_location	stadium_open	stadium_close	stadium_type	stadium_address	stadium_weather_station_code	stadium_weather_type	sta
0	Alamo Dome	San Antonio, TX	NaN	NaN	indoor	100 Montana St, San Antonio, TX 78203	78203	dome	
1	Allegiant Stadium	Paradise, NV	2020.0	NaN	indoor	3333 Al Davis Way, Las Vegas, NV 89118	NaN	dome	
2	Alltel Stadium	Jacksonville, FL	NaN	NaN	outdoor	1 TIAA Bank Field Dr, Jacksonville, FL 32202	NaN	NaN	
3	Alumni Stadium	Chestnut Hill, MA	NaN	NaN	outdoor	Perimeter Rd, Chestnut Hill, MA 02467	2467	cold	
4	Anaheim Stadium	Anaheim, CA	1980.0	1994.0	outdoor	2000 E Gene Autry Way, Anaheim, CA 92806	92806	warm	

```
In [15]: nfl_stadiums.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 106 entries, 0 to 105
Data columns (total 15 columns):
 stadium_name                106 non-null object
 stadium_location            106 non-null object
 stadium_open                82 non-null float64
 stadium_close               41 non-null float64
 stadium_type                106 non-null object
 stadium_address             106 non-null object
 stadium_weather_station_code 93 non-null object
 stadium_weather_type        99 non-null object
 stadium_capacity            45 non-null object
 stadium_surface             59 non-null object
 STATION                     55 non-null object
 NAME                        55 non-null object
 LATITUDE                    55 non-null float64
 LONGITUDE                   55 non-null float64
 ELEVATION                    55 non-null float64
dtypes: float64(5), object(10)
memory usage: 12.5+ KB
```

Now let's update all latitudes, longitudes and elevation

```
In [16]: locator = Nominatim(user_agent='myGeocoder')
```

```
In [17]: geocode = RateLimiter(locator.geocode, min_delay_seconds=1)
nfl_stadiums['location'] = nfl_stadiums['stadium_address'].apply(geocode)
nfl_stadiums['point'] = nfl_stadiums['location'].apply(lambda loc: tuple(loc.point) if loc else None)
nfl_stadiums[['latitude', 'longitude', 'elevation']] = pd.DataFrame(nfl_stadiums['point'].tolist(), index=nfl_stadiums.index)
```

```
In [18]: nfl_stadiums.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 106 entries, 0 to 105
Data columns (total 20 columns):
 stadium_name                106 non-null object
 stadium_location            106 non-null object
 stadium_open                82 non-null float64
 stadium_close               41 non-null float64
 stadium_type                106 non-null object
 stadium_address             106 non-null object
 stadium_weather_station_code 93 non-null object
 stadium_weather_type        99 non-null object
 stadium_capacity            45 non-null object
 stadium_surface             59 non-null object
 STATION                     55 non-null object
 NAME                        55 non-null object
 LATITUDE                    55 non-null float64
 LONGITUDE                   55 non-null float64
 ELEVATION                   55 non-null float64
 location                    82 non-null object
 point                       82 non-null object
 latitude                    82 non-null float64
 longitude                   82 non-null float64
 elevation                   82 non-null float64
dtypes: float64(8), object(12)
memory usage: 16.7+ KB
```

```
In [19]: nfl_stadiums.head(50)
```

44	Liberty Bowl Memorial Stadium	Memphis, TN	1997.0	1997.0	outdoor	335 S Hollywood St, Memphis, TN 38104	38104	moderate
45	Lincoln Financial Field	Philadelphia, PA	2003.0	NaN	outdoor	1020 Pattison Ave, Philadelphia, PA 19148	19148	cold
46	Los Angeles Memorial Coliseum	Los Angeles, CA	1946.0	NaN	outdoor	3911 S Figueroa St, Los Angeles, CA 90037	90037	warm

- So we didn't get all of the lats and longs we wanted, but we certainly have more than before, so we'll keep them in for now and see what happens after t
- But for the lats and longs we did get, they look to be accurate to the existing data

- Elevations all 0, so that is dissapointing, but we will move on

Remove unwanted columns

```
In [20]: nfl_stadiums = nfl_stadiums.drop(columns=['stadium_location', 'stadium_open', 'stadium_close', 'stadium_weather_type'])
nfl_stadiums.head()
```

```
Out[20]:
```

	stadium_name	stadium_type	stadium_address	stadium_weather_station_code	STATION	NAME	ELEVATION	latitude	longitude
0	Alamo Dome	indoor	100 Montana St, San Antonio, TX 78203	78203	NaN	NaN	NaN	29.416892	-98.4788
1	Allegiant Stadium	indoor	3333 Al Davis Way, Las Vegas, NV 89118	NaN	NaN	NaN	NaN	NaN	NaN
2	Alltel Stadium	outdoor	1 TIAA Bank Field Dr, Jacksonville, FL 32202	NaN	NaN	NaN	NaN	NaN	NaN
3	Alumni Stadium	outdoor	Perimeter Rd, Chestnut Hill, MA 02467	2467	NaN	NaN	NaN	NaN	NaN
4	Anaheim Stadium	outdoor	2000 E Gene Autry Way, Anaheim, CA 92806	92806	NaN	NaN	NaN	33.799711	-117.8893

```
In [ ]:
```

Reasoning for drops

- Location is worse than address
- open and close really don't matter at all for this
- Weather type is very arbitrary
- capacity doesn't really matter for what we are trying to do
- Surface has too many nulls, and there isn't enough time to look up each stadium
- LATITUDE and LONGITUDE were redunant with the ones we created
- elevation was useless, all 0

```
In [21]: nfl_stadiums.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 106 entries, 0 to 105
Data columns (total 9 columns):
 stadium_name           106 non-null object
 stadium_type           106 non-null object
 stadium_address         106 non-null object
 stadium_weather_station_code  93 non-null object
 STATION                55 non-null object
 NAME                   55 non-null object
 ELEVATION              55 non-null float64
 latitude              82 non-null float64
 longitude              82 non-null float64
dtypes: float64(3), object(6)
memory usage: 7.6+ KB
```

Move on to nfl_scores data


```
In [22]: nfl_scores.info()
Data columns (total 17 columns):
schedule_date      12947 non-null object
schedule_season    12947 non-null int64
schedule_week      12947 non-null object
schedule_playoff    12947 non-null bool
team_home          12947 non-null object
score_home         12946 non-null float64
score_away         12946 non-null float64
team_away          12947 non-null object
team_favorite_id    10468 non-null object
spread_favorite     10468 non-null float64
over_under_line     10458 non-null object
stadium            12947 non-null object
stadium_neutral    12947 non-null bool
weather_temperature 12008 non-null float64
weather_wind_mph    12008 non-null float64
weather_humidity    8388 non-null object
weather_detail      2711 non-null object
dtypes: bool(2), float64(5), int64(1), object(9)
memory usage: 1.5+ MB
```

```
In [23]: #Make schedule a datetime
nfl_scores['schedule_date'] = pd.to_datetime(nfl_scores['schedule_date'])
```

- Considering this whole project is about trying to get a classifier that can pick overs/unders we need all that data we can, any without that data will be rer

```
In [24]: nfl_scores = nfl_scores[nfl_scores['spread_favorite'].notna()]
nfl_scores = nfl_scores[nfl_scores['over_under_line'].notna()]
```

```
In [25]: nfl_scores.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10458 entries, 350 to 12946
Data columns (total 17 columns):
schedule_date      10458 non-null datetime64[ns]
schedule_season    10458 non-null int64
schedule_week      10458 non-null object
schedule_playoff    10458 non-null bool
team_home          10458 non-null object
score_home         10457 non-null float64
score_away         10457 non-null float64
team_away          10458 non-null object
team_favorite_id    10458 non-null object
spread_favorite     10458 non-null float64
over_under_line     10458 non-null object
stadium            10458 non-null object
stadium_neutral    10458 non-null bool
weather_temperature 9749 non-null float64
weather_wind_mph    9749 non-null float64
weather_humidity    6289 non-null object
weather_detail      2553 non-null object
dtypes: bool(2), datetime64[ns](1), float64(5), int64(1), object(8)
memory usage: 1.3+ MB
```

```
In [26]: nfl_scores.head(100)
```

```
Out[26]:
```

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
350	1968-01-14	1967	Superbowl	True	Green Bay Packers	33.0	14.0	Oakland Raiders	GB	-13.5
538	1969-01-12	1968	Superbowl	True	Baltimore Colts	7.0	16.0	New York Jets	IND	-18.0
727	1970-01-11	1969	Superbowl	True	Kansas City Chiefs	23.0	7.0	Minnesota Vikings	MIN	-12.0
916	1971-01-17	1970	Superbowl	True	Baltimore Colts	16.0	13.0	Dallas Cowboys	IND	-2.5
1105	1972-01-16	1971	Superbowl	True	Dallas Cowboys	24.0	3.0	Miami Dolphins	DAL	-6.0
...
2584	1979-10-08	1979	6	False	Oakland Raiders	13.0	3.0	Miami Dolphins	MIA	-1.0
2585	1979-10-14	1979	7	False	Baltimore Colts	16.0	28.0	Houston Oilers	TEN	-6.0
2586	1979-10-14	1979	7	False	Chicago Bears	7.0	27.0	New England Patriots	NE	-4.0
2587	1979-10-14	1979	7	False	Cincinnati Bengals	34.0	10.0	Pittsburgh Steelers	PIT	-10.0
2588	1979-10-14	1979	7	False	Cleveland Browns	9.0	13.0	Washington Redskins	CLE	-4.0

100 rows × 17 columns

- Only the superbowls before 1979 have the betting data, so we will remove those

```
In [27]: nfl_scores.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10458 entries, 350 to 12946
Data columns (total 17 columns):
schedule_date      10458 non-null datetime64[ns]
schedule_season    10458 non-null int64
schedule_week      10458 non-null object
schedule_playoff    10458 non-null bool
team_home          10458 non-null object
score_home         10457 non-null float64
score_away         10457 non-null float64
team_away          10458 non-null object
team_favorite_id    10458 non-null object
spread_favorite     10458 non-null float64
over_under_line     10458 non-null object
stadium            10458 non-null object
stadium_neutral     10458 non-null bool
weather_temperature 9749 non-null float64
weather_wind_mph    9749 non-null float64
weather_humidity    6289 non-null object
```

```
In [28]: nfl_scores[nfl_scores['score_away'].isnull()]
```

```
Out[28]:
```

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
12946	2021-02-07	2020	Superbowl	True	Tampa Bay Buccaneers	NaN	NaN	Kansas City Chiefs	KC	-3.

- The only data without a score_home and score_away is the last superbowl, since this game happened this will be easy enough to input

```
In [29]: nfl_scores.loc[nfl_scores['score_home'].isnull(), 'score_home'] = 31.0
nfl_scores.loc[nfl_scores['score_away'].isnull(), 'score_away'] = 9.0
```

```
In [30]: nfl_scores[nfl_scores['score_away'].isnull()]
```

```
Out[30]:
```

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite	over
--	---------------	-----------------	---------------	------------------	-----------	------------	------------	-----------	------------------	-----------------	------

```
In [31]: nfl_scores.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10458 entries, 350 to 12946
Data columns (total 17 columns):
schedule_date      10458 non-null datetime64[ns]
schedule_season    10458 non-null int64
schedule_week      10458 non-null object
schedule_playoff    10458 non-null bool
team_home          10458 non-null object
score_home         10458 non-null float64
score_away         10458 non-null float64
team_away          10458 non-null object
team_favorite_id    10458 non-null object
spread_favorite     10458 non-null float64
over_under_line     10458 non-null object
stadium            10458 non-null object
stadium_neutral     10458 non-null bool
weather_temperature 9749 non-null float64
weather_wind_mph    9749 non-null float64
weather_humidity    6289 non-null object
weather_detail      2553 non-null object
dtypes: bool(2), datetime64[ns](1), float64(5), int64(1), object(8)
memory usage: 1.3+ MB
```

- Weather detail is mostly unfilled, and mostly arbitrary, we're just gonna drop that
- I may have a different idea for adding it back in

```
In [32]: nfl_scores = nfl_scores.drop(columns='weather_detail')
```

```
In [33]: nfl_scores.info()

Int64Index: 10458 entries, 350 to 12946
Data columns (total 16 columns):
schedule_date      10458 non-null datetime64[ns]
schedule_season    10458 non-null int64
schedule_week      10458 non-null object
schedule_playoff    10458 non-null bool
team_home          10458 non-null object
score_home         10458 non-null float64
score_away         10458 non-null float64
team_away          10458 non-null object
team_favorite_id    10458 non-null object
spread_favorite     10458 non-null float64
over_under_line     10458 non-null object
stadium            10458 non-null object
stadium_neutral     10458 non-null bool
weather_temperature 9749 non-null float64
weather_wind_mph    9749 non-null float64
weather_humidity    6289 non-null object
dtypes: bool(2), datetime64[ns](1), float64(5), int64(1), object(7)
memory usage: 1.2+ MB
```

- As far as the rest of the nulls, I want to wait until we have merged with the stadium data, I have an idea that a lot of that has to do with being in an indoor
- I have concerns about over_under_line being an object

```
In [34]: nfl_scores[nfl_scores['over_under_line'].str.contains(' ')]
```

Out[34]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
2725	1979-12-23	1979	Wildcard	True	Houston Oilers	13.0	7.0	Denver Broncos	TEN	-7.0
2726	1979-12-23	1979	Wildcard	True	Philadelphia Eagles	27.0	17.0	Chicago Bears	PHI	-6.5
2727	1979-12-29	1979	Division	True	San Diego Chargers	14.0	17.0	Houston Oilers	LAC	-8.0
2728	1979-12-29	1979	Division	True	Tampa Bay Buccaneers	24.0	17.0	Philadelphia Eagles	PHI	-4.5
2729	1979-12-30	1979	Division	True	Dallas Cowboys	19.0	21.0	Los Angeles Rams	DAL	-8.5
...
4034	1986-01-04	1985	Division	True	Miami Dolphins	24.0	21.0	Cleveland Browns	MIA	-10.5
...

- Looks like we have 62 rows of data in the over_under_lines that are strings, for simplicity I am simply going to remove them

```
In [35]: nfl_scores = nfl_scores[~nfl_scores['over_under_line'].str.contains(' ')]
nfl_scores.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10396 entries, 350 to 12946
Data columns (total 16 columns):
schedule_date      10396 non-null datetime64[ns]
schedule_season    10396 non-null int64
schedule_week      10396 non-null object
schedule_playoff    10396 non-null bool
team_home          10396 non-null object
score_home         10396 non-null float64
score_away         10396 non-null float64
team_away          10396 non-null object
team_favorite_id    10396 non-null object
spread_favorite     10396 non-null float64
over_under_line     10396 non-null object
stadium            10396 non-null object
stadium_neutral    10396 non-null bool
weather_temperature 9745 non-null float64
weather_wind_mph    9745 non-null float64
weather_humidity    6289 non-null object
dtypes: bool(2), datetime64[ns](1), float64(5), int64(1), object(7)
```

```
In [36]: convert = {'over_under_line': 'float'}
nfl_scores = nfl_scores.astype(convert)
```

```
In [37]: nfl_scores.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10396 entries, 350 to 12946
Data columns (total 16 columns):
schedule_date      10396 non-null datetime64[ns]
schedule_season    10396 non-null int64
schedule_week      10396 non-null object
schedule_playoff    10396 non-null bool
team_home          10396 non-null object
score_home         10396 non-null float64
score_away         10396 non-null float64
team_away          10396 non-null object
team_favorite_id    10396 non-null object
spread_favorite     10396 non-null float64
over_under_line     10396 non-null float64
stadium            10396 non-null object
stadium_neutral     10396 non-null bool
weather_temperature 9745 non-null float64
weather_wind_mph    9745 non-null float64
weather_humidity    6289 non-null object
dtypes: bool(2), datetime64[ns](1), float64(6), int64(1), object(6)
memory usage: 1.2+ MB
```

```
In [38]: nfl_scores.describe()
```

Out[38]:

	schedule_season	score_home	score_away	spread_favorite	over_under_line	weather_temperature	weather_wind_mph
count	10396.000000	10396.000000	10396.000000	10396.000000	10396.000000	9745.000000	9745.000000
mean	2000.489708	22.707291	20.034340	-5.377982	42.093565	59.894202	7.268958
std	11.961715	10.374221	10.082921	3.431925	4.777090	15.411954	5.719802
min	1967.000000	0.000000	0.000000	-26.500000	28.000000	-6.000000	0.000000
25%	1990.000000	16.000000	13.000000	-7.000000	38.500000	50.000000	1.000000
50%	2001.000000	23.000000	20.000000	-4.500000	42.000000	64.000000	7.000000
75%	2011.000000	30.000000	27.000000	-3.000000	45.000000	72.000000	11.000000
max	2020.000000	62.000000	59.000000	0.000000	63.500000	97.000000	40.000000

Begin Merge of nfl_stadiums and nfl_scores

```
In [39]: # Need to make stadium_name equivalent to what it is (stadium) for the nfl_scores df
nfl_stadiums = nfl_stadiums.rename(columns={'stadium_name': 'stadium'})
nfl_stadiums.head()
```

Out[39]:

	stadium	stadium_type	stadium_address	stadium_weather_station_code	STATION	NAME	ELEVATION	latitude	longitude
0	Alamo Dome	indoor	100 Montana St, San Antonio, TX 78203	78203	NaN	NaN	NaN	29.416892	-98.4788
1	Allegiant Stadium	indoor	3333 Al Davis Way, Las Vegas, NV 89118	NaN	NaN	NaN	NaN	NaN	NaN
2	Alltel Stadium	outdoor	1 TIAA Bank Field Dr, Jacksonville, FL 32202	NaN	NaN	NaN	NaN	NaN	NaN
3	Alumni Stadium	outdoor	Perimeter Rd, Chestnut Hill, MA 02467	2467	NaN	NaN	NaN	NaN	NaN
4	Anaheim Stadium	outdoor	2000 E Gene Autry Way, Anaheim, CA 92806	92806	NaN	NaN	NaN	33.799711	-117.8893

```
In [40]: nfl = pd.merge(nfl_scores, nfl_stadiums, on='stadium', how='left')
nfl.head(50)
```

6	1974-01-13	1973	Superbowl	True	Miami Dolphins	24.0	7.0	Minnesota Vikings	MIA	-6.5
7	1975-01-12	1974	Superbowl	True	Minnesota Vikings	6.0	16.0	Pittsburgh Steelers	PIT	-3.0
8	1976-01-18	1975	Superbowl	True	Dallas Cowboys	17.0	21.0	Pittsburgh Steelers	PIT	-7.0
9	1977-01-09	1976	Superbowl	True	Minnesota Vikings	14.0	32.0	Oakland Raiders	LVR	-4.0
10	1978-01-15	1977	Superbowl	True	Dallas Cowboys	27.0	10.0	Denver Broncos	DAL	-6.0

```
In [41]: nfl.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10396 entries, 0 to 10395
Data columns (total 24 columns):
schedule_date      10396 non-null datetime64[ns]
schedule_season    10396 non-null int64
schedule_week      10396 non-null object
schedule_playoff    10396 non-null bool
team_home           10396 non-null object
score_home          10396 non-null float64
score_away          10396 non-null float64
team_away           10396 non-null object
team_favorite_id    10396 non-null object
spread_favorite     10396 non-null float64
over_under_line     10396 non-null float64
stadium             10396 non-null object
stadium_neutral     10396 non-null bool
weather_temperature 9745 non-null float64
weather_wind_mph    9745 non-null float64
weather_humidity    6289 non-null object
```

Evaluate those pesky stadiums without addresses/types

```
In [42]: #Seperate the areas of data without info se we can get a easier look at them
null_stad = nfl[nfl['stadium_address'].isnull()]
```

```
In [43]: null_stad['stadium'].value_counts()
```

```
Out[43]: FedEx Field                196
TIAA Bank Field                    21
Tottenham Hotspur Stadium          2
Tottenham Stadium                  1
Name: stadium, dtype: int64
```

Mostly looks like it effects a few stadiums, we should be able to add in the addresses and types pretty easily

```
In [44]: null_stad.loc[null_stad['stadium'] == 'FedEx Field', 'stadium_address'] = '1600 Fedex Way, Landover, MD 20785'
null_stad.loc[null_stad['stadium'] == 'Tottenham Hotspur Stadium', 'stadium'] = 'Tottenham Stadium'
null_stad.loc[null_stad['stadium'] == 'Tottenham Stadium', 'stadium_address'] = '782 High Rd, Tottenham, London N17 0
null_stad.loc[null_stad['stadium'] == 'TIAA Bank Field', 'stadium_address'] = '410 Franklin St, Jacksonville, FL 3220
null_stad.head(50)
```

```
/Users/matthewnykaza/opt/anaconda3/envs/learn-env/lib/python3.6/site-packages/pandas/core/indexing.py:494: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self.obj[item] = s
```



```
In [45]: null_stad.loc[null_stad['stadium'] == 'Tottenham Stadium']
```

```
Out[45]:
```

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
9682	2018-10-14	2018	6	False	Oakland Raiders	3.0	27.0	Seattle Seahawks	SEA	-3.0
9932	2019-10-06	2019	5	False	Oakland Raiders	24.0	21.0	Chicago Bears	CHI	-6.5
9950	2019-10-13	2019	6	False	Tampa Bay Buccaneers	26.0	37.0	Carolina Panthers	CAR	-2.0

3 rows × 24 columns

Let's try that geopy lat and long software again to see if we can get this thing right!

```
In [46]: null_stad['location'] = null_stad['stadium_address'].apply(geocode)
null_stad['point'] = null_stad['location'].apply(lambda loc: tuple(loc.point) if loc else None)
null_stad[['latitude', 'longitude', 'elevation']] = pd.DataFrame(null_stad['point'].tolist(), index=null_stad.index)

/Users/matthewnykaza/opt/anaconda3/envs/learn-env/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyError:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
s.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
    """Entry point for launching an IPython kernel.

/Users/matthewnykaza/opt/anaconda3/envs/learn-env/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyError:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
s.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

/Users/matthewnykaza/opt/anaconda3/envs/learn-env/lib/python3.6/site-packages/pandas/core/frame.py:3509: SettingWithCopyError:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
s.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
    self[k1] = value[k2]
```

```
In [47]: null_stad.drop(columns=['elevation', 'point', 'location'], inplace=True)
null_stad.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 220 entries, 4083 to 10385
Data columns (total 24 columns):
schedule_date          220 non-null datetime64[ns]
schedule_season        220 non-null int64
schedule_week          220 non-null object
schedule_playoff        220 non-null bool
team_home              220 non-null object
score_home             220 non-null float64
score_away             220 non-null float64
team_away              220 non-null object
team_favorite_id        220 non-null object
spread_favorite         220 non-null float64
over_under_line         220 non-null float64
stadium                220 non-null object
stadium_neutral         220 non-null bool
weather_temperature     184 non-null float64
weather_wind_mph        184 non-null float64
weather_humidity        136 non-null object
stadium_type            0 non-null object
stadium_address         220 non-null object
stadium_weather_station_code 0 non-null object
STATION                 0 non-null object
NAME                    0 non-null object
ELEVATION               0 non-null float64
latitude                220 non-null float64
longitude               220 non-null float64
dtypes: bool(2), datetime64[ns](1), float64(9), int64(1), object(11)
memory usage: 40.0+ KB
```

/Users/matthewnykaza/opt/anaconda3/envs/learn-env/lib/python3.6/site-packages/pandas/core/frame.py:4117: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
errors=errors,

Now join the dataframes back together

```
In [48]: len(nfl.columns)
```

```
Out[48]: 24
```

```
In [49]: len(null_stad.columns)
```

```
Out[49]: 24
```

```
In [50]: new_nfl = nfl.combine_first(null_stad)
```

```
In [51]: new_nfl.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10396 entries, 0 to 10395
Data columns (total 24 columns):
schedule_date          10396 non-null datetime64[ns]
schedule_season        10396 non-null int64
schedule_week          10396 non-null object
schedule_playoff       10396 non-null bool
team_home              10396 non-null object
score_home             10396 non-null float64
score_away             10396 non-null float64
team_away              10396 non-null object
team_favorite_id       10396 non-null object
spread_favorite        10396 non-null float64
over_under_line        10396 non-null float64
stadium               10396 non-null object
stadium_neutral        10396 non-null bool
weather_temperature    9745 non-null float64
weather_wind_mph       9745 non-null float64
weather_humidity       6289 non-null object
stadium_type           10176 non-null object
stadium_address        10396 non-null object
stadium_weather_station_code 10110 non-null object
STATION               7954 non-null object
NAME                   7954 non-null object
ELEVATION              7954 non-null float64
latitude               9033 non-null float64
longitude              9033 non-null float64
dtypes: bool(2), datetime64[ns](1), float64(9), int64(1), object(11)
memory usage: 1.8+ MB
```

```
In [52]: nfl.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10396 entries, 0 to 10395
Data columns (total 24 columns):
schedule_date          10396 non-null datetime64[ns]
schedule_season        10396 non-null int64
schedule_week          10396 non-null object
schedule_playoff       10396 non-null bool
team_home              10396 non-null object
score_home             10396 non-null float64
score_away             10396 non-null float64
team_away              10396 non-null object
team_favorite_id       10396 non-null object
spread_favorite        10396 non-null float64
over_under_line        10396 non-null float64
stadium                10396 non-null object
stadium_neutral        10396 non-null bool
weather_temperature    9745 non-null float64
weather_wind_mph       9745 non-null float64
weather_humidity       6289 non-null object
stadium_type           10176 non-null object
stadium_address        10176 non-null object
stadium_weather_station_code 10110 non-null object
STATION                7954 non-null object
NAME                   7954 non-null object
ELEVATION              7954 non-null float64
latitude               8813 non-null float64
longitude              8813 non-null float64
dtypes: bool(2), datetime64[ns](1), float64(9), int64(1), object(11)
memory usage: 2.2+ MB
```

```
In [53]: pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', 10)
```

```
In [54]: new_nfl[new_nfl['latitude'].isnull()]
```

...
10354	2020-12-26	2020	16	False	Las Vegas Raiders	25.0	26.0	Miami Dolphins	MIA	-2.
10361	2020-12-27	2020	16	False	Los Angeles Chargers	19.0	16.0	Denver Broncos	LAC	-2.
10369	2021-01-03	2020	17	False	Chicago Bears	16.0	35.0	Green Bay Packers	GB	-4.
10375	2021-01-03	2020	17	False	Indianapolis Colts	28.0	14.0	Jacksonville Jaguars	IND	-15.
10377	2021-01-03	2020	17	False	Los Angeles Rams	18.0	7.0	Arizona Cardinals	LAR	-1.

```
In [55]: new_nfl.loc[new_nfl['stadium'] == 'Soldier Field', 'latitude'] = 41.8623
new_nfl.loc[new_nfl['stadium'] == 'Soldier Field', 'longitude'] = -87.67167
new_nfl.loc[new_nfl['stadium_address'] == '100 S Capitol Ave, Indianapolis, IN 46225', 'longitude'] = -86.164062
new_nfl.loc[new_nfl['stadium_address'] == '100 S Capitol Ave, Indianapolis, IN 46225', 'latitude'] = 39.764705
new_nfl.loc[new_nfl['stadium_address'] == '347 Don Shula Dr, Miami Gardens, FL 33056', 'latitude'] = 25.957564
new_nfl.loc[new_nfl['stadium_address'] == '347 Don Shula Dr, Miami Gardens, FL 33056', 'longitude'] = -80.238302
new_nfl.loc[new_nfl['stadium_address'] == '500 S Capitol Ave, Indianapolis, IN 46225', 'longitude'] = -86.164062
new_nfl.loc[new_nfl['stadium_address'] == '500 S Capitol Ave, Indianapolis, IN 46225', 'latitude'] = 39.764705
new_nfl.loc[new_nfl['stadium_address'] == '1 Georgia Dome Dr, Atlanta, GA 30313', 'latitude'] = 33.757577
new_nfl.loc[new_nfl['stadium_address'] == '1 Georgia Dome Dr, Atlanta, GA 30313', 'longitude'] = -84.400952
new_nfl.loc[new_nfl['stadium_address'] == '1 Everbank Field Dr, Jacksonville, FL 32202', 'longitude'] = -81.637963
new_nfl.loc[new_nfl['stadium_address'] == '1 Everbank Field Dr, Jacksonville, FL 32202', 'latitude'] = 30.322143
new_nfl.loc[new_nfl['stadium_address'] == '1 Avenue of Champions, Clemson, SC 29634', 'latitude'] = 34.679326
new_nfl.loc[new_nfl['stadium_address'] == '1 Avenue of Champions, Clemson, SC 29634', 'longitude'] = -82.844591
new_nfl.loc[new_nfl['stadium_address'] == '1 TIAA Bank Field Dr, Jacksonville, FL 32202', 'latitude'] = 30.328008652
new_nfl.loc[new_nfl['stadium_address'] == '1 TIAA Bank Field Dr, Jacksonville, FL 32202', 'longitude'] = -81.65515899
new_nfl.loc[new_nfl['stadium_weather_station_code'] == 'Mexico City, MX', 'latitude'] = 19.303062439
new_nfl.loc[new_nfl['stadium_weather_station_code'] == 'Mexico City, MX', 'longitude'] = -99.150215149
new_nfl.loc[new_nfl['stadium_address'] == '3333 Al Davis Way, Las Vegas, NV 89118', 'latitude'] = 36.089813
new_nfl.loc[new_nfl['stadium_address'] == '3333 Al Davis Way, Las Vegas, NV 89118', 'longitude'] = -115.183925000
new_nfl.loc[new_nfl['stadium'] == 'SoFi Stadium', 'latitude'] = 33.949903
new_nfl.loc[new_nfl['stadium'] == 'SoFi Stadium', 'longitude'] = -118.343304
new_nfl.loc[new_nfl['stadium'] == 'Tulane Stadium', 'latitude'] = 29.9429822
new_nfl.loc[new_nfl['stadium'] == 'Tulane Stadium', 'longitude'] = -90.1175732
new_nfl.loc[new_nfl['stadium'] == 'Rice Stadium', 'latitude'] = 29.7163407
new_nfl.loc[new_nfl['stadium'] == 'Rice Stadium', 'longitude'] = -95.4096618
```

```
In [56]: new_nfl.info()
score_away      10396 non-null float64
team_away        10396 non-null object
team_favorite_id 10396 non-null object
spread_favorite  10396 non-null float64
over_under_line  10396 non-null float64
stadium          10396 non-null object
stadium_neutral  10396 non-null bool
weather_temperature 9745 non-null float64
weather_wind_mph    9745 non-null float64
weather_humidity    6289 non-null object
stadium_type       10176 non-null object
stadium_address    10396 non-null object
stadium_weather_station_code 10110 non-null object
STATION           7954 non-null object
NAME              7954 non-null object
ELEVATION         7954 non-null float64
latitude          10396 non-null float64
longitude         10396 non-null float64
dtypes: bool(2), datetime64[ns](1), float64(9), int64(1), object(11)
memory usage: 1.8+ MB
```

```
In [57]: new_nfl[new_nfl['stadium_type'].isnull()]
```

```
Out[57]:
```

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
4083	1997-09-14	1997	3	False	Washington Redskins	19.0	13.0	Arizona Cardinals	WAS	-8.0
4109	1997-09-28	1997	5	False	Washington Redskins	24.0	12.0	Jacksonville Jaguars	WAS	-1.0
4135	1997-10-13	1997	7	False	Washington Redskins	21.0	16.0	Dallas Cowboys	DAL	-3.0
4160	1997-10-26	1997	9	False	Washington Redskins	17.0	20.0	Baltimore Ravens	WAS	-6.0
4190	1997-11-09	1997	11	False	Washington Redskins	30.0	7.0	Detroit Lions	WAS	-4.0

```
In [58]: new_nfl.loc[new_nfl['stadium'] == 'FedEx Field', 'stadium_type'] = 'outdoor'
new_nfl.loc[new_nfl['stadium'] == 'TIAA Bank Field', 'stadium_type'] = 'outdoor'
new_nfl.loc[new_nfl['stadium'] == 'Tottenham Stadium', 'stadium_type'] = 'outdoor'
new_nfl.loc[new_nfl['stadium'] == 'Tottenham Hotspur Stadium', 'stadium_type'] = 'outdoor'
```

```
In [59]: new_nfl.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10396 entries, 0 to 10395
Data columns (total 24 columns):
schedule_date      10396 non-null datetime64[ns]
schedule_season    10396 non-null int64
schedule_week      10396 non-null object
schedule_playoff    10396 non-null bool
team_home          10396 non-null object
score_home         10396 non-null float64
score_away         10396 non-null float64
team_away          10396 non-null object
team_favorite_id    10396 non-null object
spread_favorite     10396 non-null float64
over_under_line     10396 non-null float64
stadium            10396 non-null object
stadium_neutral     10396 non-null bool
weather_temperature 9745 non-null float64
weather_wind_mph    9745 non-null float64
weather_humidity    6289 non-null object
stadium_type        10396 non-null object
stadium_address     10396 non-null object
stadium_weather_station_code 10110 non-null object
STATION            7954 non-null object
NAME               7954 non-null object
ELEVATION          7954 non-null float64
latitude           10396 non-null float64
longitude           10396 non-null float64
dtypes: bool(2), datetime64[ns](1), float64(9), int64(1), object(11)
memory usage: 1.8+ MB
```

NaNs left

- weather_temperature and weather_humidity we will allow the SimpleImputer we employ later to handle that as the mean values of both are within normal range found on the internet
- We do need to make sure that weather_humidity is in a float to make sure it can be appropriately handled by our metrics
- weather_wind we will make sure is 0 for indoor stadiums, as the mean would not be valid here. For the rest the mean will be fine
- Not worried about stadium_weather_station_code, STATION, NAME, and ELEVATION as we will drop these. We were hoping to get real weather data, but since the result these are not needed.

```
In [60]: new_nfl.loc[(new_nfl.stadium_type == 'indoor'), 'weather_wind_mph'] = 0
```

```
In [61]: new_nfl[new_nfl['weather_wind_mph'].isnull()]
```

```
Out[61]:
```

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
236	1980-01-20	1979	Superbowl	True	Los Angeles Rams	19.0	31.0	Pittsburgh Steelers	PIT	-10.0
813	1983-01-30	1982	Superbowl	True	Miami Dolphins	17.0	27.0	Washington Redskins	MIA	-3.0
1038	1984-01-22	1983	Superbowl	True	Washington Redskins	9.0	38.0	Los Angeles Raiders	WAS	-3.0
1263	1985-01-20	1984	Superbowl	True	San Francisco 49ers	38.0	16.0	Miami Dolphins	SF	-3.0
1713	1986-12-28	1986	Wildcard	True	New York Jets	35.0	15.0	Kansas City Chiefs	NYJ	-3.0
...
10390	2021-01-16	2020	Division	True	Green Bay Packers	32.0	18.0	Los Angeles Rams	GB	-7.0
10391	2021-01-17	2020	Division	True	Kansas City Chiefs	22.0	17.0	Cleveland Browns	KC	-8.0
10393	2021-01-24	2020	Conference	True	Green Bay Packers	26.0	31.0	Tampa Bay Buccaneers	GB	-3.0
10394	2021-01-24	2020	Conference	True	Kansas City Chiefs	38.0	24.0	Buffalo Bills	KC	-3.0
10395	2021-02-07	2020	Superbowl	True	Tampa Bay Buccaneers	31.0	9.0	Kansas City Chiefs	KC	-3.0

641 rows × 24 columns

```
In [62]: convert2 = {'weather_humidity': 'float'}
new_nfl = new_nfl.astype(convert2)
```

```
In [63]: new_nfl[new_nfl.weather_humidity.isnull()]
```

10391	2021-01-17	2020	Division	True	Kansas City Chiefs	22.0	17.0	Cleveland Browns	KC	-8.
10392	2021-01-17	2020	Division	True	New Orleans Saints	20.0	30.0	Tampa Bay Buccaneers	NO	-2.
10393	2021-01-24	2020	Conference	True	Green Bay Packers	26.0	31.0	Tampa Bay Buccaneers	GB	-3.
10394	2021-01-24	2020	Conference	True	Kansas City Chiefs	38.0	24.0	Buffalo Bills	KC	-3.
10395	2021-02-07	2020	Superbowl	True	Tampa Bay Buccaneers	31.0	9.0	Kansas City Chiefs	KC	-3.

4107 rows × 24 columns

```
In [64]: new_nfl.loc[(new_nfl.stadium_type == 'indoor'), 'weather_humidity'] = 50
```

```
In [65]: new_nfl.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10396 entries, 0 to 10395
Data columns (total 24 columns):
schedule_date          10396 non-null datetime64[ns]
schedule_season        10396 non-null int64
schedule_week          10396 non-null object
schedule_playoff       10396 non-null bool
team_home              10396 non-null object
score_home             10396 non-null float64
score_away             10396 non-null float64
team_away              10396 non-null object
team_favorite_id       10396 non-null object
spread_favorite        10396 non-null float64
over_under_line        10396 non-null float64
stadium                10396 non-null object
stadium_neutral        10396 non-null bool
weather_temperature    9745 non-null float64
weather_wind_mph       9755 non-null float64
weather_humidity       8197 non-null float64
stadium_weather_station_code 10396 non-null object
```

```
In [66]: new_nfl.drop(columns=['STATION', 'stadium_weather_station_code', 'NAME', 'ELEVATION'], inplace=True)
```



```
In [67]: new_nfl.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10396 entries, 0 to 10395
Data columns (total 20 columns):
schedule_date      10396 non-null datetime64[ns]
schedule_season    10396 non-null int64
schedule_week      10396 non-null object
schedule_playoff    10396 non-null bool
team_home          10396 non-null object
score_home         10396 non-null float64
score_away         10396 non-null float64
team_away          10396 non-null object
team_favorite_id    10396 non-null object
spread_favorite     10396 non-null float64
over_under_line     10396 non-null float64
stadium            10396 non-null object
stadium_neutral     10396 non-null bool
weather_temperature 9745 non-null float64
weather_wind_mph    9755 non-null float64
weather_humidity    8197 non-null float64
stadium_type        10396 non-null object
stadium_address     10396 non-null object
latitude           10396 non-null float64
longitude           10396 non-null float64
dtypes: bool(2), datetime64[ns](1), float64(9), int64(1), object(7)
memory usage: 1.5+ MB
```

Code our target

Our target variable is weather or not the over was reached, an over will be coded as a 1 and a under/push will be coded as a 0

We are also going to work towards adding division and conference information to our dataset, as these may be quite useless for our model later on

```
In [68]: # rename team_home so we can get a merge goign with the nfl_teams Dataframe
new_nfl.rename(columns={'team_home': 'team_name'}, inplace=True)
new_nfl.head()
```

Out[68]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_name	score_home	score_away	team_away	team_favorite_id	spread_favorite	over
0	1968-01-14	1967	Superbowl	True	Green Bay Packers	33.0	14.0	Oakland Raiders	GB	-13.5	
1	1969-01-12	1968	Superbowl	True	Baltimore Colts	7.0	16.0	New York Jets	IND	-18.0	
2	1970-01-11	1969	Superbowl	True	Kansas City Chiefs	23.0	7.0	Minnesota Vikings	MIN	-12.0	
3	1971-01-17	1970	Superbowl	True	Baltimore Colts	16.0	13.0	Dallas Cowboys	IND	-2.5	
4	1972-01-16	1971	Superbowl	True	Dallas Cowboys	24.0	3.0	Miami Dolphins	DAL	-6.0	

```
In [69]: # view and delete unnecessary columns in nfl_teams
nfl_teams.head()
nfl_teams.drop(columns=['team_name_short', 'team_id', 'team_id_pfr'], inplace=True)
```

```
In [70]: nfl_teams.head()
```

```
Out[70]:
```

	team_name	team_conference	team_division	team_conference_pre2002	team_division_pre2002
0	Arizona Cardinals	NFC	NFC West	NFC	NFC West
1	Phoenix Cardinals	NFC	NaN	NFC	NFC East
2	St. Louis Cardinals	NFC	NaN	NFC	NFC East
3	Atlanta Falcons	NFC	NFC South	NFC	NFC West
4	Baltimore Ravens	AFC	AFC North	AFC	AFC Central

```
In [71]: # Merge so we can get the team_division and conference information into this dataframe
nfl_2 = pd.merge(new_nfl, nfl_teams, how='left', on=['team_name'])
```

```
In [72]: nfl_2.head(20)
```

Out[72]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_name	score_home	score_away	team_away	team_favorite_id	spread_favorite
0	1968-01-14	1967	Superbowl	True	Green Bay Packers	33.0	14.0	Oakland Raiders	GB	-13.5
1	1969-01-12	1968	Superbowl	True	Baltimore Colts	7.0	16.0	New York Jets	IND	-18.0
2	1970-01-11	1969	Superbowl	True	Kansas City Chiefs	23.0	7.0	Minnesota Vikings	MIN	-12.0
3	1971-01-17	1970	Superbowl	True	Baltimore Colts	16.0	13.0	Dallas Cowboys	IND	-2.5
4	1972-01-16	1971	Superbowl	True	Dallas Cowboys	24.0	3.0	Miami Dolphins	DAL	-6.0
...
15	1979-09-02	1979	1	False	Denver Broncos	10.0	0.0	Cincinnati Bengals	DEN	-3.0
16	1979-09-02	1979	1	False	Kansas City Chiefs	14.0	0.0	Baltimore Colts	KC	-1.0
17	1979-09-02	1979	1	False	Los Angeles Rams	17.0	24.0	Oakland Raiders	LAR	-4.0
18	1979-09-02	1979	1	False	Minnesota Vikings	28.0	22.0	San Francisco 49ers	MIN	-7.0
19	1979-09-02	1979	1	False	New Orleans Saints	34.0	40.0	Atlanta Falcons	NO	-5.0

20 rows × 24 columns

```
In [73]: #Now rename appropriate home team columns in this new dataframe
nfl_2.rename(columns={'team_name': 'team_home', 'team_conference': 'team_home_conference', 'team_division': 'team_home_division'})
nfl_2.head()
```

Out[73]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite	o
0	1968-01-14	1967	Superbowl	True	Green Bay Packers	33.0	14.0	Oakland Raiders	GB	-13.5	
1	1969-01-12	1968	Superbowl	True	Baltimore Colts	7.0	16.0	New York Jets	IND	-18.0	
2	1970-01-11	1969	Superbowl	True	Kansas City Chiefs	23.0	7.0	Minnesota Vikings	MIN	-12.0	
3	1971-01-17	1970	Superbowl	True	Baltimore Colts	16.0	13.0	Dallas Cowboys	IND	-2.5	
4	1972-01-16	1971	Superbowl	True	Dallas Cowboys	24.0	3.0	Miami Dolphins	DAL	-6.0	

```
In [74]: #Do same with away teams
nfl_2.rename(columns={'team_away': 'team_name'}, inplace=True)
nfl_3 = pd.merge(nfl_2, nfl_teams, how='left', on=['team_name'])
nfl_3.rename(columns={'team_name': 'team_away', 'team_conference': 'team_away_conference', 'team_division': 'team_away_division'})
nfl_3.head()
```

Out[74]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite	o
0	1968-01-14	1967	Superbowl	True	Green Bay Packers	33.0	14.0	Oakland Raiders	GB	-13.5	
1	1969-01-12	1968	Superbowl	True	Baltimore Colts	7.0	16.0	New York Jets	IND	-18.0	
2	1970-01-11	1969	Superbowl	True	Kansas City Chiefs	23.0	7.0	Minnesota Vikings	MIN	-12.0	
3	1971-01-17	1970	Superbowl	True	Baltimore Colts	16.0	13.0	Dallas Cowboys	IND	-2.5	
4	1972-01-16	1971	Superbowl	True	Dallas Cowboys	24.0	3.0	Miami Dolphins	DAL	-6.0	

```
In [75]: # Split DF into pre-2002 and post-2002 for the sake of correctly adding team division and conference
pre_2002 = nfl_3[(nfl_3['schedule_season'] < 2002)]
post_2002 = nfl_3[(nfl_3['schedule_season'] >= 2002)]

post_2002.head()
```

Out[75]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
5321	2002-09-05	2002	1	False	New York Giants	13.0	16.0	San Francisco 49ers	SF	-4.0
5322	2002-09-08	2002	1	False	Buffalo Bills	31.0	37.0	New York Jets	NYJ	-3.0
5323	2002-09-08	2002	1	False	Carolina Panthers	10.0	7.0	Baltimore Ravens	PICK	0.0
5324	2002-09-08	2002	1	False	Chicago Bears	27.0	23.0	Minnesota Vikings	CHI	-4.5
5325	2002-09-08	2002	1	False	Cincinnati Bengals	6.0	34.0	San Diego Chargers	CIN	-3.0

```
In [76]: pre_2002.drop(columns=['team_home_conference', 'team_home_division', 'team_away_conference', 'team_away_division'], inplace=True)
pre_2002.head()
```

/Users/matthewnykaza/opt/anaconda3/envs/learn-env/lib/python3.6/site-packages/pandas/core/frame.py:4117: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
errors=errors,

Out[76]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
0	1968-01-14	1967	Superbowl	True	Green Bay Packers	33.0	14.0	Oakland Raiders	GB	-13.5
1	1969-01-12	1968	Superbowl	True	Baltimore Colts	7.0	16.0	New York Jets	IND	-18.0
2	1970-01-11	1969	Superbowl	True	Kansas City Chiefs	23.0	7.0	Minnesota Vikings	MIN	-12.0
					Baltimore Colts			Dallas Cowboys		

```
In [77]: pre_2002.rename(columns={'team_home_conference_pre2002': 'team_home_conference', 'team_home_division_pre2002': 'team_home_division'})
pre_2002.head()
```

/Users/matthewnykaza/opt/anaconda3/envs/learn-env/lib/python3.6/site-packages/pandas/core/frame.py:4238: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 return super().rename(**kwargs)

Out[77]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite	o
0	1968-01-14	1967	Superbowl	True	Green Bay Packers	33.0	14.0	Oakland Raiders	GB	-13.5	
1	1969-01-12	1968	Superbowl	True	Baltimore Colts	7.0	16.0	New York Jets	IND	-18.0	
2	1970-01-11	1969	Superbowl	True	Kansas City Chiefs	23.0	7.0	Minnesota Vikings	MIN	-12.0	
3	1971-01-17	1970	Superbowl	True	Baltimore Colts	16.0	13.0	Dallas Cowboys	IND	-2.5	
4	1972-01-16	1971	Superbowl	True	Dallas Cowboys	24.0	3.0	Miami Dolphins	DAL	-6.0	

```
In [78]: #Do the same for post 2001
post_2002.head()
```

Out[78]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite	
5321	2002-09-05	2002	1	False	New York Giants	13.0	16.0	San Francisco 49ers	SF	-4.0	
5322	2002-09-08	2002	1	False	Buffalo Bills	31.0	37.0	New York Jets	NYJ	-3.0	
5323	2002-09-08	2002	1	False	Carolina Panthers	10.0	7.0	Baltimore Ravens	PICK	0.0	
5324	2002-09-08	2002	1	False	Chicago Bears	27.0	23.0	Minnesota Vikings	CHI	-4.5	
5325	2002-09-08	2002	1	False	Cincinnati Bengals	6.0	34.0	San Diego Chargers	CIN	-3.0	

```
In [79]: post_2002.drop(columns=['team_home_conference_pre2002', 'team_home_division_pre2002', 'team_away_conference_pre2002',
post_2002.head()
```

/Users/matthewnykaza/opt/anaconda3/envs/learn-env/lib/python3.6/site-packages/pandas/core/frame.py:4117: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
errors=errors,

Out[79]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
5321	2002-09-05	2002	1	False	New York Giants	13.0	16.0	San Francisco 49ers	SF	-4.0
5322	2002-09-08	2002	1	False	Buffalo Bills	31.0	37.0	New York Jets	NYJ	-3.0
5323	2002-09-08	2002	1	False	Carolina Panthers	10.0	7.0	Baltimore Ravens	PICK	0.0
5324	2002-09-08	2002	1	False	Chicago Bears	27.0	23.0	Minnesota Vikings	CHI	-4.5
5325	2002-09-08	2002	1	False	Cincinnati Bengals	6.0	34.0	San Diego Chargers	CIN	-3.0

```
In [80]: #Put pre and post back together
all_nfl = pre_2002.append(post_2002)
all_nfl.head()
```

Out[80]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
0	1968-01-14	1967	Superbowl	True	Green Bay Packers	33.0	14.0	Oakland Raiders	GB	-13.5
1	1969-01-12	1968	Superbowl	True	Baltimore Colts	7.0	16.0	New York Jets	IND	-18.0
2	1970-01-11	1969	Superbowl	True	Kansas City Chiefs	23.0	7.0	Minnesota Vikings	MIN	-12.0
3	1971-01-17	1970	Superbowl	True	Baltimore Colts	16.0	13.0	Dallas Cowboys	IND	-2.5
4	1972-01-16	1971	Superbowl	True	Dallas Cowboys	24.0	3.0	Miami Dolphins	DAL	-6.0

```
In [81]: #Look to have a few missing divisions, but overall I'm not too concerned about that, just gonna make them 'unknown' a
all_nfl.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10396 entries, 0 to 10395
Data columns (total 24 columns):
schedule_date      10396 non-null datetime64[ns]
schedule_season    10396 non-null int64
schedule_week      10396 non-null object
schedule_playoff    10396 non-null bool
team_home          10396 non-null object
score_home         10396 non-null float64
score_away         10396 non-null float64
team_away          10396 non-null object
team_favorite_id    10396 non-null object
spread_favorite     10396 non-null float64
over_under_line     10396 non-null float64
stadium            10396 non-null object
stadium_neutral     10396 non-null bool
weather_temperature 9745 non-null float64
weather_wind_mph    9755 non-null float64
weather_humidity    8197 non-null float64
stadium_type        10396 non-null object
stadium_address     10396 non-null object
latitude           10396 non-null float64
longitude           10396 non-null float64
team_home_conference 10396 non-null object
team_home_division  10283 non-null object
team_away_conference 10396 non-null object
team_away_division  10282 non-null object
dtypes: bool(2), datetime64[ns](1), float64(9), int64(1), object(11)
memory usage: 1.8+ MB
```



```
In [82]: all_nfl.loc[all_nfl['team_home_division'].isnull(), 'team_home_division'] = 'unknown'
all_nfl.loc[all_nfl['team_away_division'].isnull(), 'team_away_division'] = 'unknown'
all_nfl.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10396 entries, 0 to 10395
Data columns (total 24 columns):
schedule_date      10396 non-null datetime64[ns]
schedule_season    10396 non-null int64
schedule_week      10396 non-null object
schedule_playoff    10396 non-null bool
team_home          10396 non-null object
score_home         10396 non-null float64
score_away         10396 non-null float64
team_away          10396 non-null object
team_favorite_id    10396 non-null object
spread_favorite     10396 non-null float64
over_under_line     10396 non-null float64
stadium            10396 non-null object
stadium_neutral     10396 non-null bool
weather_temperature 9745 non-null float64
weather_wind_mph    9755 non-null float64
weather_humidity    8197 non-null float64
stadium_type        10396 non-null object
stadium_address     10396 non-null object
latitude            10396 non-null float64
longitude           10396 non-null float64
team_home_conference 10396 non-null object
team_home_division  10396 non-null object
team_away_conference 10396 non-null object
team_away_division  10396 non-null object
dtypes: bool(2), datetime64[ns](1), float64(9), int64(1), object(11)
memory usage: 1.8+ MB
```

```
In [83]: all_nfl.head()
```

Out[83]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite	o
0	1968-01-14	1967	Superbowl	True	Green Bay Packers	33.0	14.0	Oakland Raiders	GB	-13.5	
1	1969-01-12	1968	Superbowl	True	Baltimore Colts	7.0	16.0	New York Jets	IND	-18.0	
2	1970-01-11	1969	Superbowl	True	Kansas City Chiefs	23.0	7.0	Minnesota Vikings	MIN	-12.0	
3	1971-01-17	1970	Superbowl	True	Baltimore Colts	16.0	13.0	Dallas Cowboys	IND	-2.5	
4	1972-01-16	1971	Superbowl	True	Dallas Cowboys	24.0	3.0	Miami Dolphins	DAL	-6.0	

```
In [84]: def conf(row):
        if row['team_home_conference'] == row['team_away_conference']:
            val = 1
        else:
            val = 0
        return val
```

```
In [85]: def divi(row):
        if row['team_home_division'] == row['team_away_division']:
            val = 1
        else:
            val = 0
        return val
```

```
In [86]: all_nfl['intra_conference'] = all_nfl.apply(conf, axis=1)
```

```
In [87]: all_nfl['intra_division'] = all_nfl.apply(divi, axis=1)
```

```
In [88]: all_nfl.head(-5)
```

Out[88]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
0	1968-01-14	1967	Superbowl	True	Green Bay Packers	33.0	14.0	Oakland Raiders	GB	-13.0
1	1969-01-12	1968	Superbowl	True	Baltimore Colts	7.0	16.0	New York Jets	IND	-18.0
2	1970-01-11	1969	Superbowl	True	Kansas City Chiefs	23.0	7.0	Minnesota Vikings	MIN	-12.0
3	1971-01-17	1970	Superbowl	True	Baltimore Colts	16.0	13.0	Dallas Cowboys	IND	-2.0
4	1972-01-16	1971	Superbowl	True	Dallas Cowboys	24.0	3.0	Miami Dolphins	DAL	-6.0
...
10386	2021-01-10	2020	Wildcard	True	New Orleans Saints	21.0	9.0	Chicago Bears	NO	-11.0
10387	2021-01-10	2020	Wildcard	True	Pittsburgh Steelers	37.0	48.0	Cleveland Browns	PIT	-5.0
10388	2021-01-10	2020	Wildcard	True	Tennessee Titans	13.0	20.0	Baltimore Ravens	BAL	-3.0
10389	2021-01-16	2020	Division	True	Buffalo Bills	17.0	3.0	Baltimore Ravens	BUF	-2.0
10390	2021-01-16	2020	Division	True	Green Bay Packers	32.0	18.0	Los Angeles Rams	GB	-7.0

10391 rows × 26 columns

In []:

```
In [89]: #create total, which will be the start of creating our target
all_nfl['total'] = all_nfl['score_home'] + all_nfl['score_away']
all_nfl.head()
```

Out[89]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite	o
0	1968-01-14	1967	Superbowl	True	Green Bay Packers	33.0	14.0	Oakland Raiders	GB	-13.5	
1	1969-01-12	1968	Superbowl	True	Baltimore Colts	7.0	16.0	New York Jets	IND	-18.0	
2	1970-01-11	1969	Superbowl	True	Kansas City Chiefs	23.0	7.0	Minnesota Vikings	MIN	-12.0	
3	1971-01-17	1970	Superbowl	True	Baltimore Colts	16.0	13.0	Dallas Cowboys	IND	-2.5	
4	1972-01-16	1971	Superbowl	True	Dallas Cowboys	24.0	3.0	Miami Dolphins	DAL	-6.0	

```
In [90]: # Create Target Variable for real I promise
conditional = [
    (all_nfl['total'] > all_nfl['over_under_line']),
    (all_nfl['total'] <= all_nfl['total'])]
valuez = [1, 0]
all_nfl['hit_over'] = np.select(conditional, valuez)
all_nfl.head()
```

Out[90]:

weather_humidity	stadium_type	stadium_address	latitude	longitude	team_home_conference	team_home_division	team_away_conference	team_away_division
74.0	outdoor	1501 NW 3rd St, Miami, FL 33125	25.776346	-80.219909	NFC	NFC Central	AFC	AFC We
80.0	outdoor	1501 NW 3rd St, Miami, FL 33125	25.776346	-80.219909	AFC	AFC East	NFC	AFC Ea
84.0	outdoor	Willow St. & Audubon Blvd., New Orleans, LA 70118	29.942982	-90.117573	AFC	AFC West	NFC	NFC Centr
60.0	outdoor	1501 NW 3rd St, Miami, FL 33125	25.776346	-80.219909	AFC	AFC East	NFC	NFC Ea
40.0	outdoor	Willow St. & Audubon Blvd., New Orleans, LA 70118	29.942982	-90.117573	NFC	NFC East	AFC	AFC Ea

```
In [91]: #Looks like we have a really good mix here, not to biased on one way or the other
all_nfl['hit_over'].value_counts()
```

```
Out[91]: 0    5363
         1    5033
         Name: hit_over, dtype: int64
```

```
In [92]: #Copy dataframe so I don't ruin any data unnecesarily
roll_df = all_nfl
```

```
In [93]: #Getting a rolling average of points scored for each team in the league
roll = pd.concat(
    [
        roll_df[['schedule_date', 'team_home', 'score_home']].rename(
            columns={'team_home': 'team', 'score_home': 'score'},
        ),
        roll_df[['schedule_date', 'team_away', 'score_away']].rename(
            columns={'team_away': 'team', 'score_away': 'score'},
        ),
    ], ignore_index = True,).sort_values('schedule_date')
```

```
In [94]: team_dfs = [
        roll[
            roll['team']==team
        ].set_index('schedule_date') for team in roll['team'].unique()
    ]
```

```
In [95]: for team_df in team_dfs:
        team_df['last_5'] = team_df['score'].shift(1).rolling(window=5, min_periods=1).mean()
        print(team_df, '\n')
```

	team	score	last_5
schedule_date			
1968-01-14	Green Bay Packers	33.0	NaN
1979-09-02	Green Bay Packers	3.0	33.000000
1979-09-09	Green Bay Packers	28.0	18.000000
1979-09-16	Green Bay Packers	10.0	21.333333
1979-09-23	Green Bay Packers	21.0	18.500000
...
2020-12-19	Green Bay Packers	24.0	31.400000
2020-12-27	Green Bay Packers	40.0	31.400000
2021-01-03	Green Bay Packers	35.0	33.200000
2021-01-16	Green Bay Packers	32.0	32.000000
2021-01-24	Green Bay Packers	26.0	32.400000

[704 rows x 3 columns]

	team	score	last_5
schedule_date			
1968-01-14	Oakland Raiders	14.0	NaN
1979-09-02	Oakland Raiders	33.0	33.000000

```
In [96]: last5_df = pd.concat(team_dfs)
```

```
In [97]: last5_df.reset_index(inplace=True)
```

```
In [98]: last5_df.tail(16)
```

```
Out[98]:
```

	schedule_date	team	score	last_5
20776	2020-09-13	Las Vegas Raiders	34.0	NaN
20777	2020-09-21	Las Vegas Raiders	34.0	34.000000
20778	2020-09-27	Las Vegas Raiders	20.0	34.000000
20779	2020-10-04	Las Vegas Raiders	23.0	29.333333
20780	2020-10-11	Las Vegas Raiders	40.0	27.750000
...
20787	2020-12-06	Las Vegas Raiders	31.0	24.200000
20788	2020-12-13	Las Vegas Raiders	27.0	27.200000
20789	2020-12-17	Las Vegas Raiders	27.0	26.400000
20790	2020-12-26	Las Vegas Raiders	25.0	24.400000
20791	2021-01-03	Las Vegas Raiders	32.0	23.200000

16 rows × 4 columns

```
In [99]: roll_df.rename(columns={'team_home': 'team'}, inplace=True)
roll_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10396 entries, 0 to 10395
Data columns (total 28 columns):
 schedule_date      10396 non-null datetime64[ns]
 schedule_season    10396 non-null int64
 schedule_week      10396 non-null object
 schedule_playoff    10396 non-null bool
 team               10396 non-null object
 score_home         10396 non-null float64
 score_away         10396 non-null float64
 team_away          10396 non-null object
 team_favorite_id    10396 non-null object
 spread_favorite     10396 non-null float64
 over_under_line     10396 non-null float64
 stadium            10396 non-null object
 stadium_neutral     10396 non-null bool
 weather_temperature 9745 non-null float64
 weather_wind_mph    9755 non-null float64
 weather_humidity    8197 non-null float64
 stadium_type        10396 non-null object
 stadium_address     10396 non-null object
 latitude           10396 non-null float64
 longitude           10396 non-null float64
 team_home_conference 10396 non-null object
 team_home_division  10396 non-null object
 team_away_conference 10396 non-null object
 team_away_division  10396 non-null object
 intra_conference    10396 non-null int64
 intra_division      10396 non-null int64
 total               10396 non-null float64
 hit_over            10396 non-null int64
dtypes: bool(2), datetime64[ns](1), float64(10), int64(4), object(11)
memory usage: 2.2+ MB
```

```
In [100]: new_nfl = roll_df.merge(last5_df, how='inner', left_on=['schedule_date', 'team'], right_on=['schedule_date', 'team'])
```

```
In [101]: last5_df.loc[last5_df['team']=='Seattle Seahawks']
```

Out[101]:

	schedule_date	team	score	last_5
13860	1979-09-02	Seattle Seahawks	16.0	NaN
13861	1979-09-09	Seattle Seahawks	10.0	16.000000
13862	1979-09-16	Seattle Seahawks	27.0	13.000000
13863	1979-09-23	Seattle Seahawks	34.0	17.666667
13864	1979-09-30	Seattle Seahawks	6.0	21.750000
...
14546	2020-12-13	Seattle Seahawks	40.0	22.600000
14547	2020-12-20	Seattle Seahawks	20.0	23.800000
14548	2020-12-27	Seattle Seahawks	20.0	24.600000
14549	2021-01-03	Seattle Seahawks	26.0	23.000000
14550	2021-01-09	Seattle Seahawks	20.0	23.600000

691 rows × 4 columns

```
In [102]: new_nfl.loc[new_nfl['team']=='Seattle Seahawks']
```

120	1979-10-21	1979	8	False	Seattle Seahawks	34.0	14.0	Houston Oilers	TEN	-2.0
149	1979-11-04	1979	10	False	Seattle Seahawks	0.0	24.0	Los Angeles Rams	SEA	-3.0
...
10274	2020-11-19	2020	11	False	Seattle Seahawks	28.0	21.0	Arizona Cardinals	SEA	-3.0
10314	2020-12-06	2020	13	False	Seattle Seahawks	12.0	17.0	New York Giants	SEA	-11.0
10332	2020-12-13	2020	14	False	Seattle Seahawks	40.0	3.0	New York Jets	SEA	-16.5

Los

```
In [103]: #Checking amount of NAN data, makes sense considering anytime a team has a their first 5 games, the first 4 of whice  
#will not have data  
new_nfl.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 10396 entries, 0 to 10395  
Data columns (total 30 columns):  
schedule_date      10396 non-null datetime64[ns]  
schedule_season    10396 non-null int64  
schedule_week      10396 non-null object  
schedule_playoff   10396 non-null bool  
team               10396 non-null object  
score_home         10396 non-null float64  
score_away         10396 non-null float64  
team_away          10396 non-null object  
team_favorite_id   10396 non-null object  
spread_favorite    10396 non-null float64  
over_under_line    10396 non-null float64  
stadium            10396 non-null object  
stadium_neutral    10396 non-null bool  
weather_temperature 9745 non-null float64  
weather_wind_mph   9755 non-null float64  
weather_humidity   8197 non-null float64  
stadium_type       10396 non-null object  
stadium_address    10396 non-null object  
latitude           10396 non-null float64  
longitude          10396 non-null float64  
team_home_conference 10396 non-null object  
team_home_division  10396 non-null object  
team_away_conference 10396 non-null object  
team_away_division  10396 non-null object  
intra_conference    10396 non-null int64  
intra_division      10396 non-null int64  
total              10396 non-null float64  
hit_over           10396 non-null int64  
score              10396 non-null float64  
last_5             10377 non-null float64  
dtypes: bool(2), datetime64[ns](1), float64(12), int64(4), object(11)  
memory usage: 2.3+ MB
```



```
In [104]: new_nfl.rename(columns={'team': 'team_home', 'last_5': 'last_5_home'}, inplace = True)
new_nfl.drop(columns=['score'], inplace=True)
new_nfl.tail()
```

Out[104]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	score_home	score_away	team_away	team_favorite_id	spread_favorite
10391	2021-01-17	2020	Division	True	Kansas City Chiefs	22.0	17.0	Cleveland Browns	KC	-8.5
10392	2021-01-17	2020	Division	True	New Orleans Saints	20.0	30.0	Tampa Bay Buccaneers	NO	-2.5
10393	2021-01-24	2020	Conference	True	Green Bay Packers	26.0	31.0	Tampa Bay Buccaneers	GB	-3.5
10394	2021-01-24	2020	Conference	True	Kansas City Chiefs	38.0	24.0	Buffalo Bills	KC	-3.5
10395	2021-02-07	2020	Superbowl	True	Tampa Bay Buccaneers	31.0	9.0	Kansas City Chiefs	KC	-3.5

```
In [105]: #Now do the same thing for away teams
new_nfl.rename(columns={'team_away': 'team'}, inplace=True)
```

```
In [106]: new_nfl = new_nfl.merge(last5_df, how='inner', left_on=['schedule_date', 'team'], right_on=['schedule_date', 'team'])
```

```
In [107]: #Do not want the model to know the scores of the game before hand, that would be cheating
#Any data that includes information about the score in that game will be dropped
#Not worried about stadium address, pretty arbitrary and we have other data that represents it i.e stadium name
#We'll keep the conference info for now, but I see a world where that could be dropped
new_nfl.drop(columns=['score_home', 'score_away', 'stadium_address', 'total', 'score'], inplace=True)
```

```
In [108]: new_nfl.rename(columns={'team': 'team_away', 'last_5': 'last_5_away'}, inplace = True)
```

```
In [109]: new_nfl.tail()
```

Out[109]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	team_away	team_favorite_id	spread_favorite	over_under_line	stadium
10391	2021-01-17	2020	Division	True	Kansas City Chiefs	Cleveland Browns	KC	-8.0	56.0	Arrowhead Stadium
10392	2021-01-17	2020	Division	True	New Orleans Saints	Tampa Bay Buccaneers	NO	-2.5	53.0	Mercedes-Benz Superdome
10393	2021-01-24	2020	Conference	True	Green Bay Packers	Tampa Bay Buccaneers	GB	-3.0	53.0	Lambeau Field
10394	2021-01-24	2020	Conference	True	Kansas City Chiefs	Buffalo Bills	KC	-3.0	55.0	Arrowhead Stadium
10395	2021-02-07	2020	Superbowl	True	Tampa Bay Buccaneers	Kansas City Chiefs	KC	-3.0	56.0	Raymond James Stadium

```
In [110]: #Check work
new_nfl[(new_nfl == 'Seattle Seahawks').any(axis=1)]

/Users/matthewnykaza/opt/anaconda3/envs/learn-env/lib/python3.6/site-packages/pandas/core/ops/__init__.py:1115: FutureWarning: Comparison of DataFrame with a scalar is deprecated. In the future will perform elementwise comparison
  result = method(y)
```

```
Out[110]:
```

le_season	schedule_week	schedule_playoff	team_home	team_away	team_favorite_id	spread_favorite	over_under_line	stadium	stadium_neutral	weather_t
1979	1	False	Seattle Seahawks	San Diego Chargers	SEA	-2.0	42.5	Seattle Kingdome	False	
1979	2	False	Miami Dolphins	Seattle Seahawks	MIA	-7.0	40.5	Orange Bowl	False	
1979	3	False	Seattle Seahawks	Oakland Raiders	SEA	-3.0	44.0	Seattle Kingdome	False	
1979	4	False	Denver Broncos	Seattle Seahawks	DEN	-6.0	37.0	Mile High Stadium	False	
1979	5	False	Seattle Seahawks	Kansas City Chiefs	SEA	-6.0	43.0	Seattle Kingdome	False	
...	
2020	14	False	Seattle Seahawks	New York Jets	SEA	-16.5	49.0	CenturyLink Field	False	
2020	15	False	Washington Football Team	Seattle Seahawks	SEA	-6.0	44.0	FedEx Field	False	
2020	16	False	Seattle Seahawks	Los Angeles Rams	SEA	-1.5	48.0	CenturyLink Field	False	
2020	17	False	San Francisco 49ers	Seattle Seahawks	SEA	-7.0	45.0	Levi's Stadium	False	
2020	Wildcard	True	Seattle Seahawks	Los Angeles Rams	SEA	-3.0	42.5	CenturyLink Field	False	

```
In [142]: #Make sure schedule_week is a numerical (will take multiple steps)
new_nfl['schedule_week'] = new_nfl['schedule_week'].replace({'18': '17'})
```

```
In [148]: new_nfl['schedule_week'] = new_nfl['schedule_week'].replace({'Wildcard': '18', 'WildCard': '18', 'Division': '19',
                                                                    'Conference': '20', 'Superbowl': '21', 'SuperBowl': '21'})
```

```
In [151]: new_nfl['schedule_week'] = new_nfl['schedule_week'].astype('int32')
```

```
In [ ]:
```

```
In [192]: new_nfl.info()
team_home      10361 non-null object
team_away      10361 non-null object
team_favorite_id 10361 non-null object
spread_favorite 10361 non-null float64
over_under_line 10361 non-null float64
stadium_neutral 10361 non-null int64
weather_temperature 9712 non-null float64
weather_wind_mph 9722 non-null float64
weather_humidity 8166 non-null float64
stadium_type    10361 non-null object
latitude        10361 non-null float64
longitude       10361 non-null float64
intra_conference 10361 non-null int64
intra_division  10361 non-null int64
hit_over        10361 non-null int64
last_5_home     10361 non-null float64
last_5_away     10361 non-null float64
estimated_total  10361 non-null float64
dtypes: float64(10), int32(1), int64(6), object(4)
memory usage: 1.7+ MB
```

```
In [161]: #Create new column based on expected point total (based on last5 scores added)
new_nfl['estimated_total'] = new_nfl['last_5_home'] + new_nfl['last_5_away']
new_nfl.tail()
```

Out[161]:

under_line	stadium_neutral	weather_temperature	weather_wind_mph	weather_humidity	stadium_type	latitude	longitude	intra_conference	intra_division	hit_over
56.0	0	NaN	NaN	NaN	outdoor	39.048939	-94.483984	1		0
53.0	0	NaN	0.0	50.0	indoor	29.951049	-90.082308	1		1
53.0	0	NaN	NaN	NaN	outdoor	44.500958	-88.061034	1		0
55.0	0	NaN	NaN	NaN	outdoor	39.048939	-94.483984	1		0
56.0	0	NaN	NaN	NaN	outdoor	27.977901	-82.505322	0		0

```
In [152]: new_nfl.describe()
```

```
Out[152]:
```

	spread_favorite	over_under_line	stadium_neutral	weather_temperature	weather_wind_mph	weather_humidity	latitude	longitude	intra_conference	intra_division
count	10361.000000	10361.000000	10361.000000	9712.000000	9722.000000	8166.000000	10361.000000	10361.000000	10361.000000	10361.000000
mean	-5.378294	42.109420	0.008590	59.860379	7.257972	62.686383	37.854069	-90.307461	0.726667	0.445692
std	3.431416	4.770111	0.092287	15.417128	5.723770	15.693731	5.260936	15.843484	0.445692	0.445692
min	-26.500000	28.000000	0.000000	-6.000000	0.000000	4.000000	19.303062	-122.389227	0.000000	0.000000
25%	-7.000000	38.500000	0.000000	50.000000	1.000000	50.000000	33.757577	-95.407756	0.000000	0.000000
50%	-4.500000	42.000000	0.000000	63.000000	7.000000	62.000000	39.098319	-86.164062	1.000000	0.000000
75%	-3.000000	45.000000	0.000000	72.000000	11.000000	75.000000	41.506056	-80.014015	1.000000	1.000000
max	0.000000	63.500000	1.000000	97.000000	40.000000	100.000000	51.590914	-0.069968	1.000000	1.000000

```
In [219]: new_nfl.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10361 entries, 6 to 10395
Data columns (total 20 columns):
schedule_season      10361 non-null int64
schedule_week        10361 non-null int32
schedule_playoff      10361 non-null int64
team_home             10361 non-null object
team_away             10361 non-null object
spread_favorite       10361 non-null float64
over_under_line       10361 non-null float64
stadium_neutral       10361 non-null int64
weather_temperature   9712 non-null float64
weather_wind_mph      9722 non-null float64
weather_humidity      8166 non-null float64
stadium_type          10361 non-null object
latitude              10361 non-null float64
longitude             10361 non-null float64
intra_conference      10361 non-null int64
intra_division        10361 non-null int64
hit_over              10361 non-null int64
last_5_home           10361 non-null float64
last_5_away           10361 non-null float64
estimated_total       10361 non-null float64
dtypes: float64(10), int32(1), int64(6), object(3)
memory usage: 1.6+ MB
```

```
In [157]: #Remove categorical data that already has information associated in data (stadium address, stadium name, etc.)
new_nfl.drop(columns=['team_home_conference', 'team_home_division', 'team_away_conference', 'team_away_division',
                    'stadium', 'schedule_date', 'team_favorite_id'], inplace=True)
```

```
In [158]: #Going to drop NaN data in the last_5 categories, represents a tiny % of the data, should be fine moving forward
new_nfl.dropna(subset=['last_5_home', 'last_5_away'], inplace=True)
```

```
In [114]: hist_data = ['schedule_season', 'schedule_week', 'spread_favorite', 'over_under_line',
                    'weather_temperature', 'weather_wind_mph', 'weather_humidity', 'stadium_type',
                    'latitude', 'longitude', 'intra_conference', 'intra_division', 'hit_over', 'last_5_home', 'last_5_away']
```

Time for Modeling

```
In [115]: def histogram_view(data_set):
          """This will produce 3 columns of histograms and a corresponding number
          of rows depending on the len(data_set.columns). The result will be a
          side by side view of all of the histograms of each column in a data set"""
          ncols = 3
          nrows = int(np.ceil(len(data_set.columns) / (1.0*ncols)))
          fig, axes = plt.subplots(nrows=nrows, ncols=ncols, figsize=(20, 15))
          # Lazy counter so we can remove unwated axes
          counter = 0
          for i in range(nrows):
              for j in range(ncols):

                  ax = axes[i][j]

                  # Plot when we have data
                  if counter < len(data_set.columns):

                      ax.hist(data_set[data_set.columns[counter]], bins=20, color='blue', alpha=0.5, label='{}'.format(data_set.columns[counter]))
                      ax.set_xlabel('{}'.format(data_set.columns[counter]))
                      ax.set_ylabel('Density')
                      leg = ax.legend(loc='best')
                      leg.draw_frame(True)
                      ax.grid(which='both', axis='both', linestyle='-')
                  # Remove axis when we no longer have data
                  else:
                      ax.set_axis_off()

                  counter += 1

          plt.show()
```


```
In [116]: new_nfl.head()
```

Out[116]:

	schedule_date	schedule_season	schedule_week	schedule_playoff	team_home	team_away	team_favorite_id	spread_favorite	over_under_line	stadium
6	1974-01-13	1973	Superbowl	True	Miami Dolphins	Minnesota Vikings	MIA	-6.5	33.0	Rice Stadium
8	1976-01-18	1975	Superbowl	True	Dallas Cowboys	Pittsburgh Steelers	PIT	-7.0	36.0	Orange Bowl
9	1977-01-09	1976	Superbowl	True	Minnesota Vikings	Oakland Raiders	LVR	-4.0	38.0	Rose Bowl
11	1979-01-21	1978	Superbowl	True	Dallas Cowboys	Pittsburgh Steelers	PIT	-3.5	37.0	Orange Bowl
16	1979-09-02	1979	1	False	Kansas City Chiefs	Baltimore Colts	KC	-1.0	37.0	Arrowhead Stadium


```
In [155]: # Numpy doesn't like dtypes 'bool', so we'll just make them true = 1 and false = 0 for simplicity
          new_nfl[['schedule_playoff', 'stadium_neutral']] = new_nfl[['schedule_playoff', 'stadium_neutral']].replace({True: 1, False: 0})
```

```
In [195]: #Check for normalcy
histogram_view(new_nfl)
```



- Overall the data looks very normal and should be a great starting point for regression

```
In [348]: # Now how about multicollinearity
plt.figure(figsize=(8,8))
matrix = np.triu(new_nfl.corr())
sns.heatmap(new_nfl.corr(), mask=matrix)
plt.show()
```



There are a lot of variables that appear to have a similar level of correlation in the middle of this correlation matrix by pulling up a Variance Inflation Factor


```
In [359]: #Calculating VIF to review multicollinearity
VIF_cols = []
for c in new_nfl.columns:
    if new_nfl[c].dtype in ['float64', 'int64', 'int32']:
        VIF_cols.append(c)
vif = pd.DataFrame()
X = new_nfl[VIF_cols].dropna()
vif["variables"] = X.columns
vif["VIF"] = [variance_inflation_factor(X.values, i)
              for i in range(len(X.columns))]

print(vif)
```

	variables	VIF
0	schedule_season	335.896910
1	schedule_week	6.841824
2	schedule_playoff	1.241090
3	spread_favorite	3.910372
4	over_under_line	146.000099
..
12	intra_division	2.086455
13	hit_over	1.955315
14	last_5_home	inf
15	last_5_away	inf
16	estimated_total	inf

[17 rows x 2 columns]

```
/Users/matthewnykaza/opt/anaconda3/envs/learn-env/lib/python3.6/site-packages/statsmodels/stats/outliers_influence.py
o encountered in double_scalars
vif = 1. / (1. - r_squared_i)
```

- In general anything above 10 is considered to be high correlation, and just on a brief view it is clear that we have some extraordinarily high correlations going on
- This gives some credence to the idea we will need to use a model (i.e. Tree Based) that does not care about multicollinearity

```
In [306]: #Look into object cols
new_nfl[[c for c in new_nfl.columns if new_nfl[c].dtype == 'object']].describe()
```

Out[306]:

	team_home	team_away	stadium_type
count	10361	10361	10361
unique	42	42	3
top	New England Patriots	Green Bay Packers	outdoor
freq	362	350	7866

```
In [307]: #going to to test train splits a little differently, as we want to test on the most recent data as this is likely
# what we will be doing in the future
#We are going to used the last 5 seasons of data 2016-2020 as based ont schedule season for our test, this is a tad
#small, but we can try and change this later if we find it is too small
# and we'll define X and y in this step
train = new_nfl[(new_nfl['schedule_season'] < 2016)]
test = new_nfl[(new_nfl['schedule_season'] >= 2016)]
y_train = train['hit_over']
y_test = test['hit_over']
X_train = train.drop(columns=['hit_over'])
X_test = test.drop(columns=['hit_over'])
```

```
In [308]: #Check to make sure indicies look right for train and test
y_train.head()
```

```
Out[308]: 6      0
8      1
9      1
11     1
16     0
Name: hit_over, dtype: int64
```

```
In [309]: X_train.head()
```

```
Out[309]:
```

	schedule_season	schedule_week	schedule_playoff	team_home	team_away	spread_favorite	over_under_line	stadium_neutral	weather_temperature	weather
6	1973	21	1	Miami Dolphins	Minnesota Vikings	-6.5	33.0	1	47.0	
8	1975	21	1	Dallas Cowboys	Pittsburgh Steelers	-7.0	36.0	1	49.0	
9	1976	21	1	Minnesota Vikings	Oakland Raiders	-4.0	38.0	1	52.0	
11	1978	21	1	Dallas Cowboys	Pittsburgh Steelers	-3.5	37.0	1	71.0	
16	1979	1	0	Kansas City Chiefs	Baltimore Colts	-1.0	37.0	0	76.0	

```
In [310]: y_test.head()
```

```
Out[310]: 9059    1
9060    0
9061    1
9062    0
9063    0
Name: hit_over, dtype: int64
```

```
In [311]: X_test.head()
```

```
Out[311]:
```

	schedule_season	schedule_week	schedule_playoff	team_home	team_away	spread_favorite	over_under_line	stadium_neutral	weather_temperature	weather_humidity	weather_wind_mph
9059	2016	1	0	Denver Broncos	Carolina Panthers	-3.0	40.5	0	82.0	72.0	12.0
9060	2016	1	0	Arizona Cardinals	New England Patriots	-8.5	44.0	0	72.0	72.0	12.0
9061	2016	1	0	Atlanta Falcons	Tampa Bay Buccaneers	-2.5	47.0	0	72.0	72.0	12.0
9062	2016	1	0	Baltimore Ravens	Buffalo Bills	-3.0	44.5	0	82.0	72.0	12.0
9063	2016	1	0	Dallas Cowboys	New York Giants	-1.0	47.5	0	72.0	72.0	12.0

All looks solid for this point, time to get to modeling.

Modeling

```
In [312]: # Look at some basic information about how we might want to set up our pipeline for the following steps
```

```
num_cols = []
ohe_cols = []
freq_cols = []

for c in X_train.columns:
    if new_nfl[c].dtype in ['float64', 'int64', 'int32']:
        num_cols.append(c)
    elif len(X_train[c].unique()) <= 15:
        ohe_cols.append(c)
    else:
        freq_cols.append(c)
```

```
In [313]: # Check our work
```

```
print(f"Numeric: {num_cols}")
print(f"To OHE: {ohe_cols}")
print(f"To Frequency Encode: {freq_cols}")
```

```
Numeric: ['schedule_season', 'schedule_week', 'schedule_playoff', 'spread_favorite', 'over_under_line', 'stadium_neutral', 'stadium_type', 'er_wind_mph', 'weather_humidity', 'latitude', 'longitude', 'intra_conference', 'intra_division', 'last_5_home', 'last_5_away']
To OHE: ['stadium_type']
To Frequency Encode: ['team_home', 'team_away']
```

- I even might just want to OHE the home and away team names, but I also might just get rid of these all together as I have some concerns about how team performance is changing over time, relying on how a team performed in 1979, just to learn trends in the data
- I think it makes sense to OHE the stadium type, but if this gives me issues later I can also make them 0=outdoors, 1=indoors and 3=retractable, but I am not sure about the other two

```
In [314]: #Check vs. test to make sure there are not any issues
num_cols_test = []
ohe_cols_test = []
freq_cols_test = []

for c in X_test.columns:
    if new_nfl[c].dtype in ['float64', 'int64', 'int32']:
        num_cols_test.append(c)
    elif len(X_test[c].unique()) <= 15:
        ohe_cols_test.append(c)
    else:
        freq_cols_test.append(c)
```

```
In [315]: # Check our work
print(f"Numeric: {num_cols_test}")
print(f"To OHE: {ohe_cols_test}")
print(f"To Frequency Encode: {freq_cols_test}")
```

```
Numeric: ['schedule_season', 'schedule_week', 'schedule_playoff', 'spread_favorite', 'over_under_line', 'stadium_neutral_wind_mph', 'weather_humidity', 'latitude', 'longitude', 'intra_conference', 'intra_division', 'last_5_home', 'last_5_away']
To OHE: ['stadium_type']
To Frequency Encode: ['team_home', 'team_away']
```

Setting up preprocessing

- Numerical Data
 - Since (as seen in the histograms) there do not seem to be many major outliers present in the data I am simply going to use a Min-Max scaler to get the data to be on the same scale
 - For the imputer I am going to start with the median, there really isn't that much a difference than the mean here so I doubt that it will have much of an effect
- Categorical Data
 - Nothing too crazy here, only have the one column of stadium_type, so this should be straightforward encoding
 - May later on try the home and away teams here, but the more I think about it, the more likely it is that I will drop them.
- Frequency Data
 - As stated before I very well may drop the teams, just going to encode for now with the Count Encoder

```
In [316]: X_train.describe()
```

```
Out[316]:
```

	schedule_season	schedule_week	schedule_playoff	spread_favorite	over_under_line	stadium_neutral	weather_temperature	weather_wind_mph	weather_humidity
count	9027.000000	9027.000000	9027.000000	9027.000000	9027.000000	9027.000000	8782.000000	8782.000000	7027.000000
mean	1997.961560	9.433920	0.036889	-5.389110	41.537244	0.007422	59.496242	7.578684	65.250000
std	10.541066	5.179595	0.188501	3.426065	4.541411	0.085837	15.322526	5.731902	14.727500
min	1973.000000	1.000000	0.000000	-26.500000	28.000000	0.000000	-6.000000	0.000000	28.000000
25%	1989.000000	5.000000	0.000000	-7.000000	38.000000	0.000000	49.000000	2.000000	55.000000
50%	1998.000000	10.000000	0.000000	-4.500000	41.000000	0.000000	63.000000	8.000000	65.000000
75%	2007.000000	14.000000	0.000000	-3.000000	44.500000	0.000000	72.000000	11.000000	70.000000
max	2015.000000	21.000000	1.000000	0.000000	63.000000	1.000000	95.000000	40.000000	95.000000

```
In [317]: X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 9027 entries, 6 to 9058
Data columns (total 19 columns):
schedule_season      9027 non-null int64
schedule_week        9027 non-null int32
schedule_playoff     9027 non-null int64
team_home            9027 non-null object
team_away            9027 non-null object
spread_favorite      9027 non-null float64
over_under_line      9027 non-null float64
stadium_neutral      9027 non-null int64
weather_temperature  8782 non-null float64
weather_wind_mph     8782 non-null float64
weather_humidity     7987 non-null float64
stadium_type         9027 non-null object
latitude             9027 non-null float64
longitude            9027 non-null float64
intra_conference     9027 non-null int64
intra_division       9027 non-null int64
last_5_home          9027 non-null float64
last_5_away          9027 non-null float64
estimated_total      9027 non-null float64
dtypes: float64(10), int32(1), int64(5), object(3)
memory usage: 1.3+ MB
```

```
In [318]: # Now, set up the preprocessing steps for each type of col
```

```
num_transformer = Pipeline(steps=[
    ('num_imputer', SimpleImputer(strategy='median')),
    ('scaler', MinMaxScaler())])

ohe_transformer = Pipeline(steps=[
    ('cat_imputer', SimpleImputer(strategy='constant', fill_value='Unknown')),
    ('ohe', OneHotEncoder(handle_unknown='ignore'))])

freq_transformer = Pipeline(steps=[
    ('freq_imputer', SimpleImputer(strategy='constant', fill_value='Unknown')),
    ('freq_enc', ce.CountEncoder(normalize=True,
                                handle_unknown=0,
                                min_group_size=0.001,
                                min_group_name='Other'))])
```

```
In [319]: # Put together our preprocessor using a Column Transformer
```

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', num_transformer, num_cols),
        ('cat_ohe', ohe_transformer, ohe_cols),
        ('cat_freq', freq_transformer, freq_cols)])
```

```
In [320]: #Check work
X_train.info()
schedule_season      9027 non-null int64
schedule_week        9027 non-null int64
schedule_playoff      9027 non-null object
team_home             9027 non-null object
team_away             9027 non-null object
spread_favorite       9027 non-null float64
over_under_line       9027 non-null float64
stadium_neutral       9027 non-null int64
weather_temperature   8782 non-null float64
weather_wind_mph      8782 non-null float64
weather_humidity      7987 non-null float64
stadium_type          9027 non-null object
latitude              9027 non-null float64
longitude             9027 non-null float64
intra_conference      9027 non-null int64
intra_division        9027 non-null int64
last_5_home           9027 non-null float64
last_5_away           9027 non-null float64
estimated_total       9027 non-null float64
dtypes: float64(10), int32(1), int64(5), object(3)
memory usage: 1.3+ MB
```

```
In [538]: # Append classifier to preprocessing pipeline.
# Now we have a full prediction pipeline.
clf_logreg = Pipeline(steps=[('preprocessor', preprocessor),
                              ('classifier', LogisticRegression(class_weight='balanced', random_state=42))])

clf_logreg.fit(X_train, y_train)
```

```
Out[538]: Pipeline(steps=[('preprocessor',
                           ColumnTransformer(transformers=[('num',
                                                             Pipeline(steps=[('num_imputer',
                                                                 SimpleImputer(strategy='median')),
                                                                 ('scaler',
                                                                  MinMaxScaler()))],
                                                             ['schedule_season',
                                                              'schedule_week',
                                                              'schedule_playoff',
                                                              'spread_favorite',
                                                              'over_under_line',
                                                              'stadium_neutral',
                                                              'weather_temperature',
                                                              'weather_wind_mph',
                                                              'weather_humidity',
                                                              'latitude', 'longitude',
                                                              'intra_conference',
                                                              'intra_division',
                                                              'last_5_home', 'last_5_away',
                                                              'estimated_total']),
                                                             ('cat_ohe',
                                                              Pipeline(steps=[('cat_imputer',
                                                                 SimpleImputer(fill_value='Unknown',
                                                                 strategy='constant')),
                                                                 ('ohe',
                                                                  OneHotEncoder(handle_unknown='ignore'))]),
                                                             ['stadium_type'])])),
                           ('classifier',
                            LogisticRegression(class_weight='balanced', random_state=42))])
```

```

In [539]: # This is just a nice little bit of code that will give us relevant test results of our model!
def evaluate(estimator, X_train, X_test, y_train, y_test, use_decision_function='yes'):
    """
    Evaluation function to show a few scores for both the train and test set
    Also shows a confusion matrix for the test set

    use_decision_function allows you to toggle whether you use decision_function or
    predict_proba in order to get the output needed for roc_auc_score
    If use_decision_function == 'skip', then it ignores calculating the roc_auc_score

    Additionally for models that have a decision function this model will show a ROC Curve
    """
    # grab predictions
    train_preds = estimator.predict(X_train)
    test_preds = estimator.predict(X_test)

    # output needed for roc_auc_score
    if use_decision_function == 'skip': # skips calculating the roc_auc_score
        train_out = False
        test_out = False
    elif use_decision_function == 'yes': # not all classifiers have decision_function
        train_out = estimator.decision_function(X_train)
        test_out = estimator.decision_function(X_test)
    elif use_decision_function == 'no':
        train_out = estimator.predict_proba(X_train)[: , 1] # proba for the 1 class
        test_out = estimator.predict_proba(X_test)[: , 1]
    else:
        raise Exception ("The value for use_decision_function should be 'skip', 'yes' or 'no'.")

    print(type(test_out))

    # print scores
    print("Train Scores")
    print("-----")
    print(f"Accuracy: {accuracy_score(y_train, train_preds)}")
    print(f"F1 Score: {f1_score(y_train, train_preds)}")
    if type(train_out) == np.ndarray:
        print(f"ROC-AUC: {roc_auc_score(y_train, train_out)}")
    print("----" * 5)
    print("Test Scores")
    print("-----")
    print(f"Accuracy: {accuracy_score(y_test, test_preds)}")
    print(f"F1 Score: {f1_score(y_test, test_preds)}")
    if type(test_out) == np.ndarray:
        print(f"ROC-AUC: {roc_auc_score(y_test, test_out)}")

    # plot test confusion matrix
    plot_confusion_matrix(estimator, X_test, y_test)
    plt.show()

    #Plot ROC Curve
    if use_decision_function == 'yes':
        y_train_score = estimator.decision_function(X_train)
        y_test_score = estimator.decision_function(X_test)

        train_fpr, train_tpr, train_thresholds = roc_curve(y_train, y_train_score)
        test_fpr, test_tpr, test_thresholds = roc_curve(y_test, y_test_score)

        print('Train AUC: {}'.format(auc(train_fpr, train_tpr)))
        print('Test AUC: {}'.format(auc(test_fpr, test_tpr)))

```

```
plt.figure(figsize=(10, 8))
lw = 2

plt.plot(train_fpr, train_tpr, color='blue',
         lw=lw, label='Train ROC curve')
plt.plot(test_fpr, test_tpr, color='darkorange',
         lw=lw, label='Test ROC curve')

plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.yticks([i/20.0 for i in range(21)])
plt.xticks([i/20.0 for i in range(21)])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

else:
    None
```



```
In [541]: evaluate(clf_logreg, X_train, X_test, y_train, y_test)
```

```
<class 'numpy.ndarray'>
```

```
Train Scores
```

```
-----
```

```
Accuracy: 0.5324027916251246
```

```
F1 Score: 0.5266345183357631
```

```
ROC-AUC: 0.5416435185207935
```

```
-----
```


```
Test Scores
```

```
-----
```

```
Accuracy: 0.5112443778110944
```

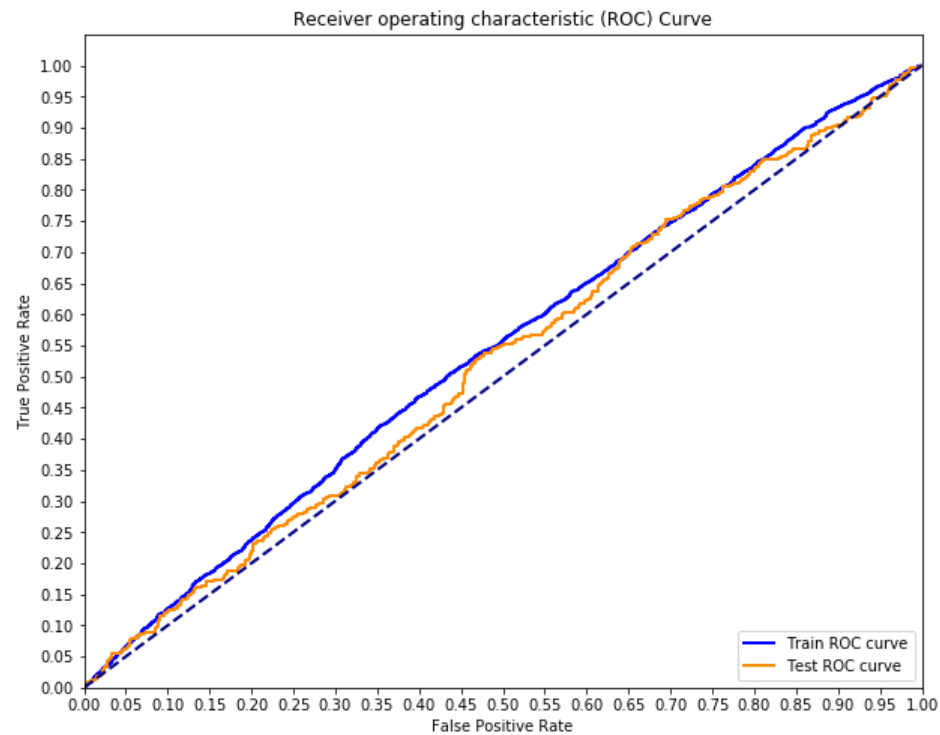
```
F1 Score: 0.44651952461799665
```

```
ROC-AUC: 0.523093815249478
```



```
Train AUC: 0.5416435185207935
```

```
Test AUC: 0.523093815249478
```



Some Takeaways

- Slightly better than a guess
- The confusion matrix shows that we are best at telling when a game is going to be 'under'
- ROC curve indicates that we are essentially guessing with this model

Let's take teams out and see what happens???

```
In [542]: #Remove team names from the data
no_team_names = new_nfl.copy().drop(columns=['team_home', 'team_away'])
train = no_team_names[(no_team_names['schedule_season'] < 2016)]
test = no_team_names[(no_team_names['schedule_season'] >= 2016)]
y_train = train['hit_over']
y_test = test['hit_over']
X_train = train.drop(columns=['hit_over'])
X_test = test.drop(columns=['hit_over'])
```

```
In [543]: #Need to run these again
num_cols = []
ohe_cols = []
freq_cols = []

for c in X_train.columns:
    if no_team_names[c].dtype in ['float64', 'int64', 'int32']:
        num_cols.append(c)
    elif len(X_train[c].unique()) <= 15:
        ohe_cols.append(c)
    else:
        freq_cols.append(c)
```

```
In [544]: preprocessor = ColumnTransformer(
    transformers=[
        ('num', num_transformer, num_cols),
        ('cat_ohe', ohe_transformer, ohe_cols)])
```

```
In [545]: print(f"Numeric: {num_cols}")
print(f"To OHE: {ohe_cols}")
print(f"To Frequency Encode: {freq_cols}")
```

```
Numeric: ['schedule_season', 'schedule_week', 'schedule_playoff', 'spread_favorite', 'over_under_line', 'stadium_neut
er_wind_mph', 'weather_humidity', 'latitude', 'longitude', 'intra_conference', 'intra_division', 'last_5_home', 'last
To OHE: ['stadium_type']
To Frequency Encode: []
```

```

In [546]: clf_logreg = Pipeline(steps=[('preprocessor', preprocessor),
                                       ('classifier', LogisticRegression(class_weight='balanced', random_state=42))])
clf_logreg.fit(X_train, y_train)

Out[546]: Pipeline(steps=[('preprocessor',
                           ColumnTransformer(transformers=[('num',
                                                            Pipeline(steps=[('num_imputer',
                                                                 SimpleImputer(strategy='median')),
                                                                 ('scaler',
                                                                  MinMaxScaler()))],
                                                            ['schedule_season',
                                                             'schedule_week',
                                                             'schedule_playoff',
                                                             'spread_favorite',
                                                             'over_under_line',
                                                             'stadium_neutral',
                                                             'weather_temperature',
                                                             'weather_wind_mph',
                                                             'weather_humidity',
                                                             'latitude', 'longitude',
                                                             'intra_conference',
                                                             'intra_division',
                                                             'last_5_home', 'last_5_away',
                                                             'estimated_total']),
                                                            ('cat_ohe',
                                                             Pipeline(steps=[('cat_imputer',
                                                                 SimpleImputer(fill_value='Unknown',
                                                                 strategy='constant')),
                                                                 ('ohe',
                                                                  OneHotEncoder(handle_unknown='ignore'))]),
                                                             ['stadium_type'])])),
                           ('classifier',
                            LogisticRegression(class_weight='balanced', random_state=42))])

```

```
In [547]: evaluate(clf_logreg, X_train, X_test, y_train, y_test)
```

```
<class 'numpy.ndarray'>
```

```
Train Scores
```

```
-----
```

```
Accuracy: 0.5324027916251246
```

```
F1 Score: 0.5266345183357631
```

```
ROC-AUC: 0.5416435185207935
```

```
-----
```


```
Test Scores
```

```
-----
```

```
Accuracy: 0.5112443778110944
```


```
F1 Score: 0.44651952461799665
```

```
ROC-AUC: 0.523093815249478
```



```
Train AUC: 0.5416435185207935
```

```
Test AUC: 0.523093815249478
```



Some Takeaways

- Overall it does not appear to have made much of a difference
- We may be ever so slightly overfit, but not any more so than what we were before
- It is clear that a basic Logistic Regression is not the answer
 - This is likely because there is a whole lot of multicollinearity in the data, so I believe it will be best to use a Tree based model as this will negate any n


Tree Based Models

```
In [596]: tree_logreg = Pipeline(steps=[('preprocessor', preprocessor),
                                         ('classifier', DecisionTreeClassifier(class_weight='balanced', random_state=42))])
tree_logreg.fit(X_train, y_train)
```

```
Out[596]: Pipeline(steps=[('preprocessor',  
    ColumnTransformer(transformers=[('num',  
        Pipeline(steps=[('num_imputer',  
            SimpleImputer(strategy='median')),  
            ('scaler',  
                MinMaxScaler())]),  
        [ 'schedule_season',  
          'schedule_week',  
          'schedule_playoff',  
          'spread_favorite',  
          'over_under_line',  
          'stadium_neutral',  
          'weather_temperature',  
          'weather_wind_mph',  
          'weather_humidity',  
          'latitude', 'longitude',  
          'intra_conference',  
          'intra_division',  
          'last_5_home', 'last_5_away',  
          'estimated_total']),  
        ('cat_ohe',  
        Pipeline(steps=[('cat_imputer',  
            SimpleImputer(fill_value='Unknown',  
                            strategy='constant')),  
            ('ohe',  
                OneHotEncoder(handle_unknown='ignore'))]),  
        [ 'stadium_type'])])),  
    ('classifier',  
    DecisionTreeClassifier(class_weight='balanced',  
                           random_state=42))])
```

```
In [597]: evaluate(tree_logreg, X_train, X_test, y_train, y_test, use_decision_function='no')
```

```
<class 'numpy.ndarray'>
Train Scores
-----
Accuracy: 1.0
F1 Score: 1.0
ROC-AUC: 1.0
-----
Test Scores
-----
Accuracy: 0.49025487256371814
F1 Score: 0.4670846394984326
ROC-AUC: 0.48937108503131976
```



Some Takeaways

- Scores are largely unchanged in the grand scheme of things here, but with some hyperparameter tuning I believe that we can get a solid score
- Current goal is to get in the 60 range for all scores

```
In [615]: #Let's use GridsearchCV
param_grid = {
    'classifier__max_depth': [x for x in range(4,12,2)],
    'classifier__min_samples_split': [x for x in range(4, 12, 2)],
    'classifier__max_features': [ 'log2'],
    'classifier__min_samples_leaf': [x for x in range(10, 20, 2)],
}
```

```
In [616]: gridpy = GridSearchCV(tree_logreg, param_grid, n_jobs=1, cv=3)
```



```
In [617]: gridy.fit(X_train, y_train)
```


```
Out[617]: GridSearchCV(cv=3,
                      estimator=Pipeline(steps=[('preprocessor',
                                                ColumnTransformer(transformers=[('num',
                                                                                      Pipeline(steps=[('num_imputer',
                                                                                          SimpleImputer(strategy='median'),
                                                                                          ('scaler',
                                                                                          MinMaxScaler()))]),
                                                                                      ['schedule_season',
                                                                                      'schedule_week',
                                                                                      'schedule_playoff',
                                                                                      'spread_favorite',
                                                                                      'over_under_line',
                                                                                      'stadium_neutral',
                                                                                      'weather_temperature',
                                                                                      'weather_wind_mph',
                                                                                      'weather_...
                                                                                      strategy='co
                                                                                      ('ohe',
                                                                                      OneHotEncoder(handle_unknown='ignore'),
                                                                                      ['stadium_type'])))]),
                                                ('classifier',
                                                DecisionTreeClassifier(class_weight='balanced',
                                                                        random_state=42))]),
                      n_jobs=1,
                      param_grid={'classifier__max_depth': [4, 6, 8, 10],
                                  'classifier__max_features': ['log2'],
                                  'classifier__min_samples_leaf': [10, 12, 14, 16, 18],
                                  'classifier__min_samples_split': [4, 6, 8, 10]})
```

```
In [618]: print(f"Best parameter's score: {gridy.best_score_:0.3f}:")
          print(gridy.best_params_)
```

```
Best parameter's score: 0.521):
{'classifier__max_depth': 8, 'classifier__max_features': 'log2', 'classifier__min_samples_leaf': 14, 'classifier__min_samples_split': 10}
```

```
In [619]: evaluate(gridify.best_estimator_, X_train, X_test, y_train, y_test, use_decision_function='no')
```

```
<class 'numpy.ndarray'>
Train Scores
-----
Accuracy: 0.5624238395923341
F1 Score: 0.47529224229543043
ROC-AUC: 0.5956151261262331
-----
Test Scores
-----
Accuracy: 0.5269865067466267
F1 Score: 0.44109831709477415
ROC-AUC: 0.5279222766217871
```



On the Hyperparameters

- criterion - from what I could find, either way the model performed about the same, so I am not concerned with this
- splitter - similar takeaway as criterion
- max_depth - with nothing set the tree will be perfect on the training set, but not generalize well on test set. Through some testing I found that a depth of this seemed a bit small to me
- min_samples_split - looked like 4 was very consistently the best number here, higher numbers tend to introduce under-fitting, so I think 4 was a good number
- min_samples_leaf - is similar to min_samples_split, but when used together they can help us make sure that the tree isn't using 1 sample to make a decision
- max_features - mostly helps with computational time, so I really went with allowing the model to determine what is best, as I am not too concerned with speed
- max_leaf_nodes - ran this a few times and it appeared that having a number around 80 was best for this metric, although none of this improved the model's performance, so I left it alone.

Some Takeaways

- The results largely improved, with the exception of F1 score on the test set, this is not a great development as the lower f1 score (combined with greater variance) is a bad sign
- Seems like this was better at selecting overs (1), but may have just erred that way too often
- It seems that through pruning I am mostly getting trees that err one way or the other, but not really improving greatly on their overall scores


One last model, Random Forests

This will be an ensemble model that ideally will be the most powerful

[illegible]

```
In [579]: evaluate(forest, X_train, X_test, y_train, y_test, use_decision_function='no')
```

```
<class 'numpy.ndarray'>
Train Scores
-----
Accuracy: 1.0
F1 Score: 1.0
ROC-AUC: 1.0
-----
Test Scores
-----
Accuracy: 0.525487256371814
F1 Score: 0.4373333333333333
ROC-AUC: 0.5209315825473395
```



Some Takeaways

- The model is picking the under (0) far too often, seemingly the opposite of our last model
- On a whole the scores did not improve much, over the DecisionTree with hyperparameter selection

```
In [599]: # Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Method of selecting samples for training each tree
bootstrap = [True, False]
```

```
In [606]: random_grid = {'classifier__n_estimators': n_estimators,
                        'classifier__max_features': max_features,
                        'classifier__max_depth': max_depth,
                        'classifier__min_samples_split': min_samples_split,
                        'classifier__min_samples_leaf': min_samples_leaf,
                        'classifier__bootstrap': bootstrap}
print(random_grid)
```

```
{'classifier__n_estimators': [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000], 'classifier__max_features': ['sqrt', 'log2', 'best'], 'classifier__min_samples_split': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110, None], 'classifier__min_samples_leaf': [5, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110, None], 'classifier__bootstrap': [True, False]}
```

```
In [607]: forest = Pipeline(steps=[('preprocessor', preprocessor),
                                   ('classifier', RandomForestClassifier(class_weight='balanced', random_state=42))])
forest.fit(X_train, y_train)
```

```
Out[607]: Pipeline(steps=[('preprocessor',  
    ColumnTransformer(transformers=[('num',  
        Pipeline(steps=[('num_imputer',  
            SimpleImputer(strategy='median')),  
            ('scaler',  
                MinMaxScaler()))],  
        [ 'schedule_season',  
          'schedule_week',  
          'schedule_playoff',  
          'spread_favorite',  
          'over_under_line',  
          'stadium_neutral',  
          'weather_temperature',  
          'weather_wind_mph',  
          'weather_humidity',  
          'latitude', 'longitude',  
          'intra_conference',  
          'intra_division',  
          'last_5_home', 'last_5_away',  
          'estimated_total']),  
      ('cat_ohe',  
        Pipeline(steps=[('cat_imputer',  
            SimpleImputer(fill_value='Unknown',  
                          strategy='constant')),  
            ('ohe',  
              OneHotEncoder(handle_unknown='ignore'))]),  
        [ 'stadium_type'])])),  
    ('classifier',  
      RandomForestClassifier(class_weight='balanced',  
                            random_state=42))])
```

```
In [609]: # Use the random grid to search for best hyperparameters  
# First create the base model to tune
```

```
# Random search of parameters, using 3 fold cross validation,  
# search across 20 different combinations, and use all available cores  
rf_random = RandomizedSearchCV(estimator = forest, param_distributions = random_grid, n_iter = 20, cv = 3, verbose=2,  
# Fit the random search model  
rf_random.fit(X_train, y_train)
```

Fitting 3 folds for each of 20 candidates, totalling 60 fits

```
Out[609]: RandomizedSearchCV(cv=3,  
                             estimator=Pipeline(steps=[('preprocessor',  
                                                         ColumnTransformer(transformers=[('num',  
                                                                 Pipeline(steps=[('num_imputer',  
                                                                 SimpleImputer(strategy=  
                                                                 ('scaler',  
                                                                 MinMaxScaler()))]),  
                                                                 ['schedule_season',  
                                                                 'schedule_week',  
                                                                 'schedule_playoff',  
                                                                 'spread_favorite',  
                                                                 'over_under_line',  
                                                                 'stadium_neutral',  
                                                                 'weather_temperature',  
                                                                 'weather_wind_mph',  
                                                                 'we...  
                                                         ],  
                                                         n_iter=20, n_jobs=-1,  
                                                         param_distributions={'classifier__bootstrap': [True, False],  
                                                                 'classifier__max_depth': [10, 20, 30,  
                                                                 40, 50, 60,  
                                                                 70, 80, 90,  
                                                                 100, 110,  
                                                                 None],  
                                                                 'classifier__max_features': ['auto',  
                                                                 'sqrt'],  
                                                                 'classifier__min_samples_leaf': [1, 2,  
                                                                 4],  
                                                                 'classifier__min_samples_split': [2, 5,  
                                                                 10],  
                                                                 'classifier__n_estimators': [200, 400,  
                                                                 600, 800,  
                                                                 1000, 1200,  
                                                                 1400, 1600,  
                                                                 1800,  
                                                                 2000]}},  
                                                         verbose=2)
```

```
In [622]: #This will give us a good baseline as to where to start our GridsearchCV process  
rf_random.best_params_
```

```
Out[622]: {'classifier__n_estimators': 200,  
           'classifier__min_samples_split': 5,  
           'classifier__min_samples_leaf': 4,  
           'classifier__max_features': 'auto',  
           'classifier__max_depth': 110,  
           'classifier__bootstrap': True}
```

```
In [626]: #Utilized those paramaters to narrow down on GridSearchCV
param_grid = {'classifier__n_estimators': [150, 200, 250],
              'classifier__max_features': ['auto'],
              'classifier__max_depth': [110, 120, 140],
              'classifier__min_samples_split': [4, 5, 6],
              'classifier__min_samples_leaf': [3, 4, 5],
              'classifier__bootstrap': [True]}
```

```
In [627]: gridy_2 = GridSearchCV(estimator=forest, param_grid=param_grid, n_jobs=1, cv=3)
```

```
In [628]: gridy_2.fit(X_train, y_train)
```


```
Out[628]: GridSearchCV(cv=3,
                      estimator=Pipeline(steps=[('preprocessor',
                                                ColumnTransformer(transformers=[('num',
                                                                                      Pipeline(steps=[('num_imputer',
                                                                                          SimpleImputer(strategy='median'),
                                                                                          ('scaler',
                                                                                          MinMaxScaler()))],
                                                                                      ['schedule_season',
                                                                                      'schedule_week',
                                                                                      'schedule_playoff',
                                                                                      'spread_favorite',
                                                                                      'over_under_line',
                                                                                      'stadium_neutral',
                                                                                      'weather_temperature',
                                                                                      'weather_wind_mph',
                                                                                      'weather_...',
                                                                                      OneHotEncoder(handle_unknown='ignore',
                                                                                      ['stadium_type'])))],
                                                ('classifier',
                                                RandomForestClassifier(class_weight='balanced',
                                                                        random_state=42))]),
                      n_jobs=1,
                      param_grid={'classifier__bootstrap': [True],
                                  'classifier__max_depth': [110, 120, 140],
                                  'classifier__max_features': ['auto'],
                                  'classifier__min_samples_leaf': [3, 4, 5],
                                  'classifier__min_samples_split': [4, 5, 6],
                                  'classifier__n_estimators': [150, 200, 250]})
```

```
In [629]: print(f"Best parameter's score: {gridy_2.best_score_:0.3f}")
print(gridy_2.best_params_)
```

```
Best parameter's score: 0.521)
{'classifier__bootstrap': True, 'classifier__max_depth': 110, 'classifier__max_features': 'auto', 'classifier__min_sam
ples_split': 4, 'classifier__n_estimators': 150}
```

```
In [630]: evaluate(grid2y_2.best_estimator_, X_train, X_test, y_train, y_test, use_decision_function='no')
```

```
<class 'numpy.ndarray'>
Train Scores
-----
Accuracy: 0.9996676636756398
F1 Score: 0.9996567112941984
ROC-AUC: 0.9999885999578788
-----
Test Scores
-----
Accuracy: 0.5352323838080959
F1 Score: 0.45035460992907805
ROC-AUC: 0.5317247462020304
```



Conclusion

- This model was able to perform the best out of all previous models, and I really think that finding RancomizerCV as a method of getting a good starting point was a good idea.
- All metrics rose by around .01 - .015 points, which may not seem like a lot, but this was the greatest increase in the data seen to date.
- Overall what this shows is the need to complete more data cleaning, and get more data.
- I believe that one major issue is that I am trying to beat Las Vegas, which creates these point spreads using models much more advanced and practiced than I am. This could be a viable product.
- It may be that I have too much past data, when they game was very different, this could be a hinderence as well

Further Work

- Need to get more information about each individual game, this could include more data mining and more feature engineering of feature variables
- More tests of Hyperparameters

- Utilize graphing techniques to help determine some of these features
- Try boost models
 - They tend to be more powerful, and may be able to achieve better scores

In []: